



**TRAILHEAD**  
TECHNOLOGY PARTNERS

# Warm and Fuzzy

Semantic Search in .NET



Jonathan “J.” Tower

# Medication Status Lookup

Please enter the name of the medication below and click on a result in the dropdown below.

The medication database does not contain information on, or that applies to any dietary ingredient.

**Search Tip:** Search for the generic name first (acetaminophen). If the medication is not found, search for the brand name (Tylenol).

Medication Name

 Medication name

Additional Information

Additional information, if any, will be displayed here.

TRAILHEAD TECHNOLOGY PARTNERS

HOME MEDICATION STATUS LOOKUP SUPPLEMENT INFORMATION ABOUT JOSH+PRO...@TRAILHEADTECHNOLOGY.COM

## Medication Status Lookup

Please enter the name of the medication below and click on a result in the dropdown below.  
The medication database does not contain information on, or that applies to any dietary ingredient.

**Search Tip:** Search for the generic name first (acetaminophen). If the medication is not found, search for the brand name (Tylenol).

Medication Name

Additional Information

Additional information, if any, will be displayed here.

tylenol ≠ acetaminophen

advil ≠ ibuprofen ≠ NSAID

[HOME](#)[\*\*MEDICATION STATUS  
LOOKUP\*\*](#)[SUPPLEMENT  
INFORMATION](#)[ABOUT](#) ▾[JOSH+PRO...@TRAILHEADTECHNOLOGY.COM](#) ▾

# Medication Status Lookup

Please enter the name of the medication below and click on a result in the dropdown below.

The medication database does not contain information on, or that applies to any dietary ingredient.

**Search Tip:** Search for the generic name first (acetaminophen). If the medication is not found, search for the brand name (Tylenol).

Medication Name

 Medication name

Additional Information

Additional information, if any, will be displayed here.

Learn how to add  
AI-powered  
**semantic search**  
to your .NET apps



# Jonathan "J." Tower

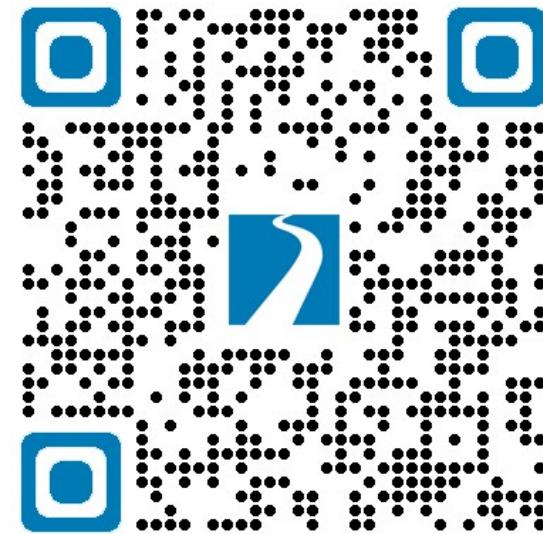
Partner & Principal Consultant



- ✉ Microsoft MVP in .NET
- ✉ .NET Foundation Board of Directors
- ✉ jtower@trailheadtechnology.com
- ✉ trailheadtechnology.com/blog
- ✉ jtowermi
- ✉ Jonathan "J." Tower

[github.com/trailheadtechnology/api-security](https://github.com/trailheadtechnology/api-security)

## EXPERT Consultation



bit.ly/th-offer

# The Evolution of Fuzzy Search

## Exact Match

```
SELECT * FROM Products WHERE Name = 'car'
```

## Exact Match

```
SELECT * FROM Products WHERE Name = 'car'
```

User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“ca” won’t match “car”

“kar” won’t match “car”

## Exact Match

```
SELECT * FROM Products WHERE Name = 'car'
```

### User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“ca” won’t match “car”

“kar” won’t match “car”

## Exact Match

```
SELECT * FROM Products WHERE Name = 'car'
```

### User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“ca” won’t match “car”

“kar” won’t match “car”

## Exact Match

```
SELECT * FROM Products WHERE Name = 'car'
```

### User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“ca” won’t match “car”

“kar” won’t match “car”

## LIKE Queries

```
SELECT * FROM Products WHERE Name LIKE '%car%'
```

## LIKE Queries

```
SELECT * FROM Products WHERE Name LIKE '%car%'
```

User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“kar” won’t match “car”

# LIKE Queries

```
SELECT * FROM Products WHERE Name LIKE '%car%'
```

## User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“kar” won’t match “car”

## LIKE Queries

```
SELECT * FROM Products WHERE Name LIKE '%car%'
```

### User frustration:

“automobile” won’t match “car”

“cra” won’t match “car”

“kar” won’t match “car”

# Levenshtein

**Levenshtein("kitten", "sitting") = 3**

L	e	v	e	n	s	h	t	e	i	n		
L	e	v	e	n	s	h	t	e	i	n		
L	e	v	e	n	s	h	t	e	i	n		
0	0.5	1	1.5	2	2.5	3	3.5	4	4.5	5	5.5	
L	0.5	0	0.5	1	1.5	2	2.5	3	3.5	4	4.5	5
e	1	0.5	0	0.5	1	1.5	2	2.5	3	3.5	4	4.5
v	1.5	1	0.5	0	0.5	1	1.5	2	2.5	3	3.5	4
e	2	1.5	1	0.5	0	0.5	1	1.5	2	2.5	3	3.5
n	2.5	2	1.5	1	0.5	0	0.5	1	1.5	2	2.5	3
s	3	2.5	2	1.5	1	0.5	0	0.5	1	1.5	2	2.5
h	3.5	3	2.5	2	1.5	1	0.5	0	0.5	1	1.5	2
t	4	3.5	3	2.5	2	1.5	1	0.5	0	0.5	1	1.5
e	4.5	4	3.5	3	2.5	2	1.5	1	0.5	0	0.5	1
i	5	4.5	4	3.5	3	2.5	2	1.5	1	0.5	0	0.5
n	5.5	5	4.5	4	3.5	3	2.5	2	1.5	1	0.5	0

**Rules (configurable):**

Substitutions cost 1

Deletion or insertion costs 1

**Ex:**

kitten → sitten: 1

sitten → sittin: 1

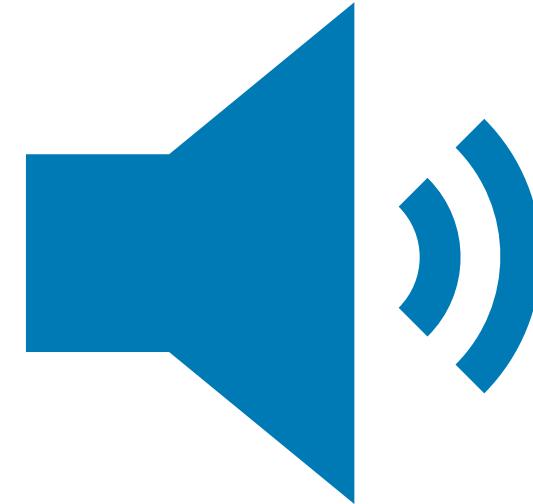
sittin → sitting: 1

Total: 3

Video: Substitution as 1 and cost of deletion or insertion as 0.5

# Soundex

Creates **4-character code** based on how they **sound**, not how they're spelled



## Rules:

- Keeps the first letter of the word
- Converts the rest into numbers representing consonant sounds
- Drops vowels and silent letters
- Words that sound similar → same code

## Ex:

"Smith" → S530  
"Smyth" → S530

"Robert" → R163  
"Rupert" → R163

# From Fuzzy to Semantic Search



# From Fuzzy to Semantic Search



## Fuzzy Search

Find things that **look similar**



## Semantic Search

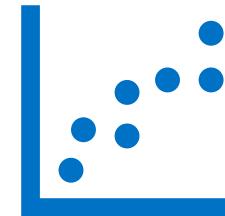
Find things that **mean the same**

# Core Concepts of Semantic Search

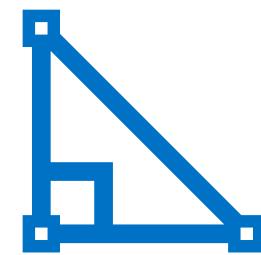
# Core Concepts of Semantic Search



Vectors



Embeddings

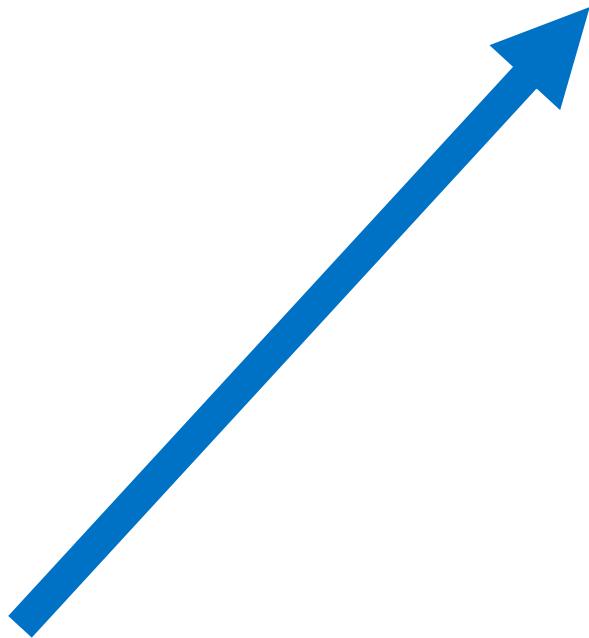


Cosine  
Differences

# Vectors

# Vectors

1. Direction
2. Magnitude

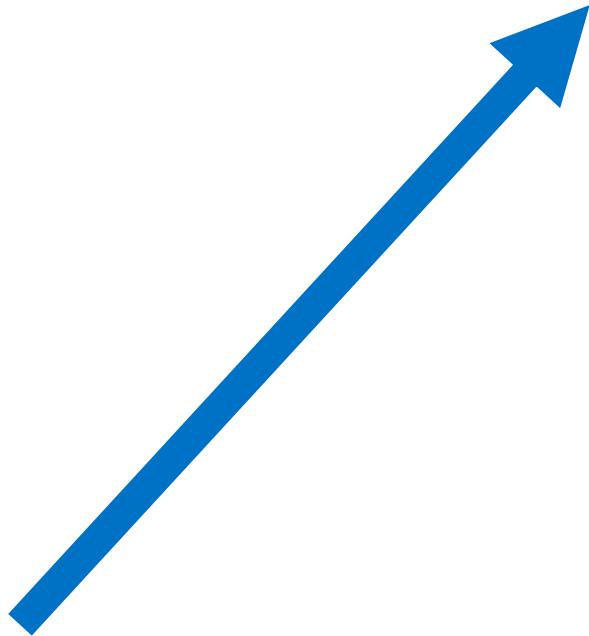


# Vectors

1. Direction
2. Magnitude

**Ex:**

1 mile northwest



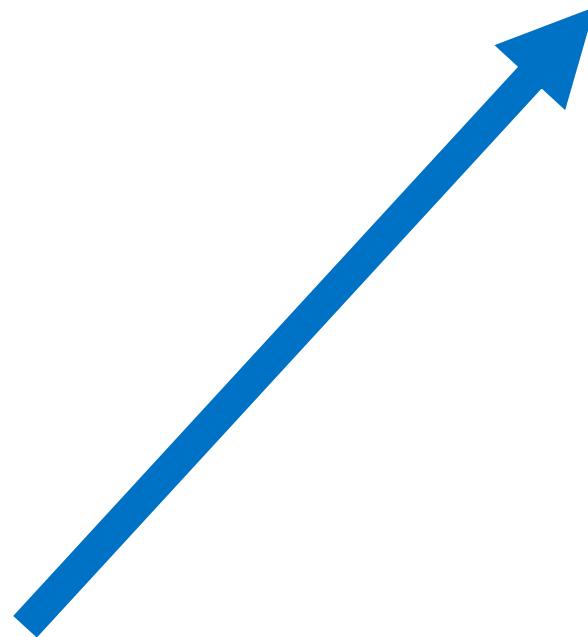
# Vectors

1. Direction
2. Magnitude

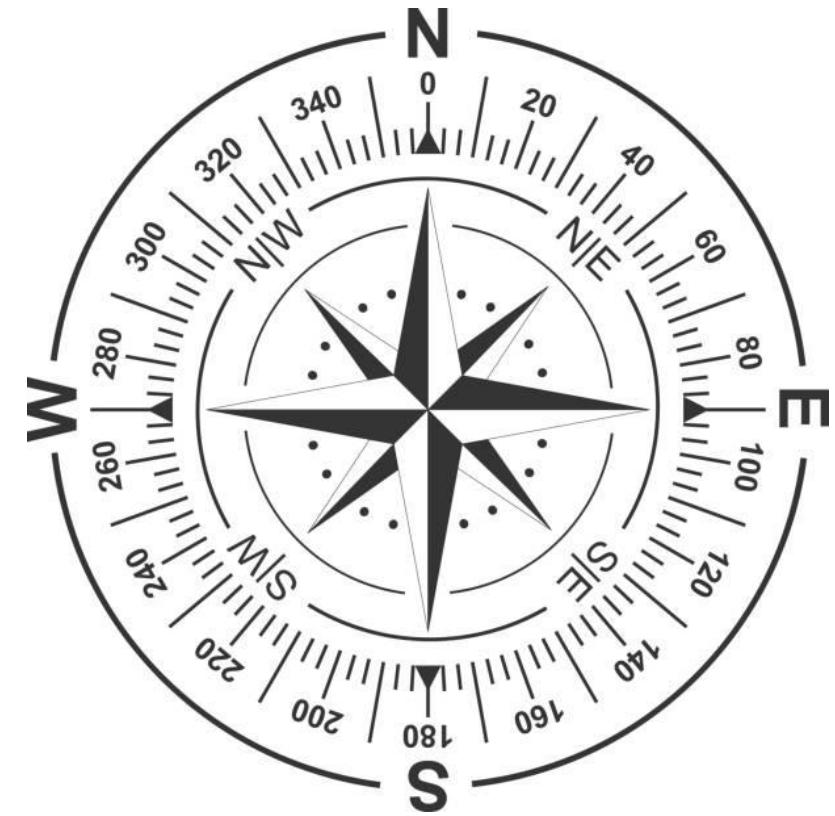
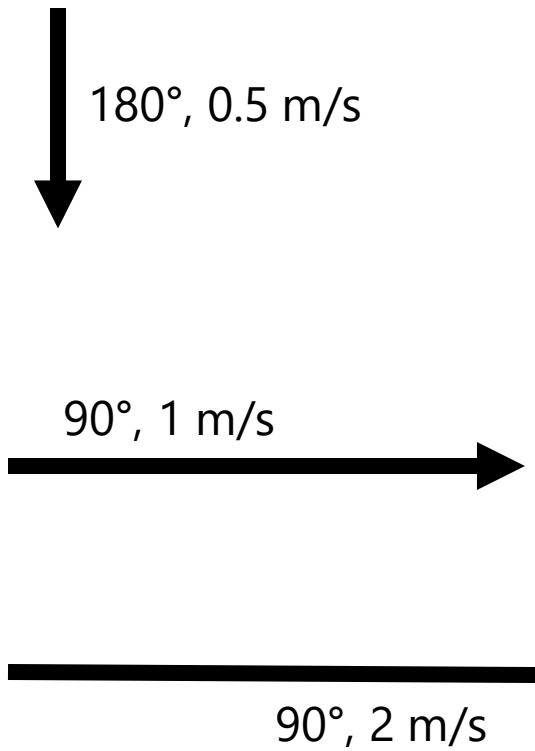
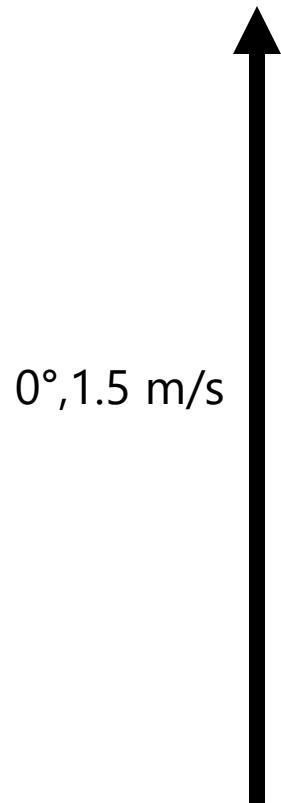
**Ex:**

1 mile northwest

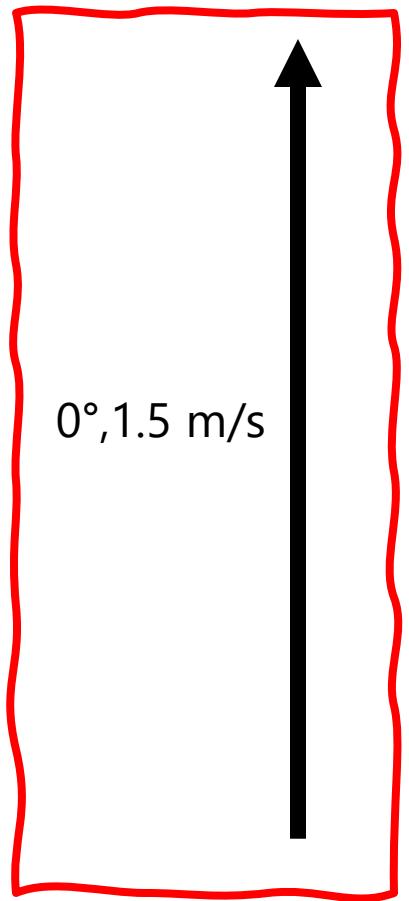
1 m/s at  $130^\circ$



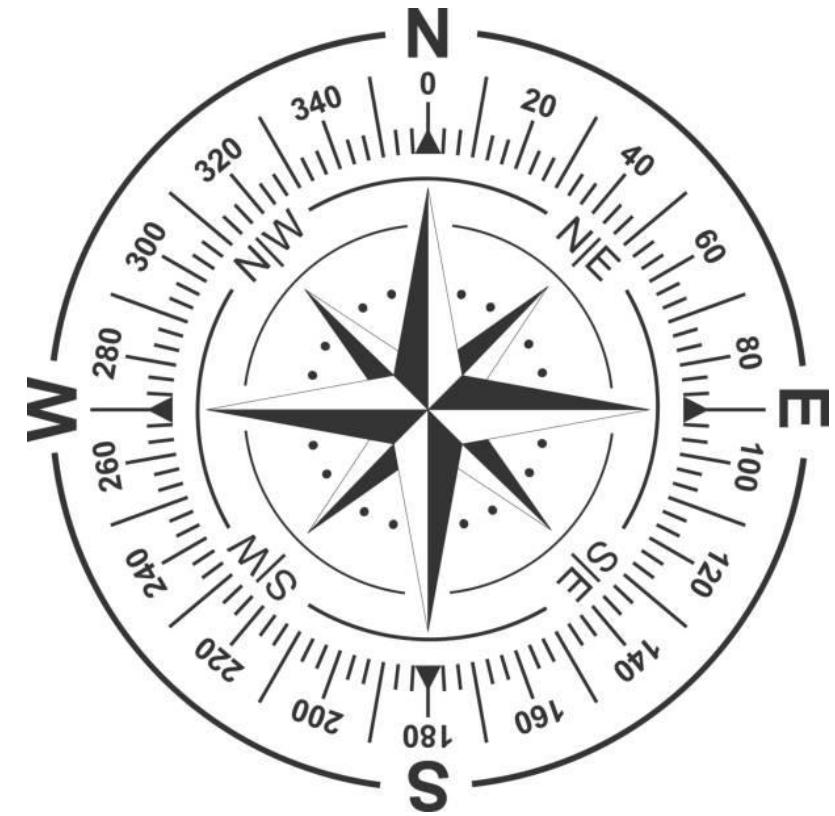
# Vectors



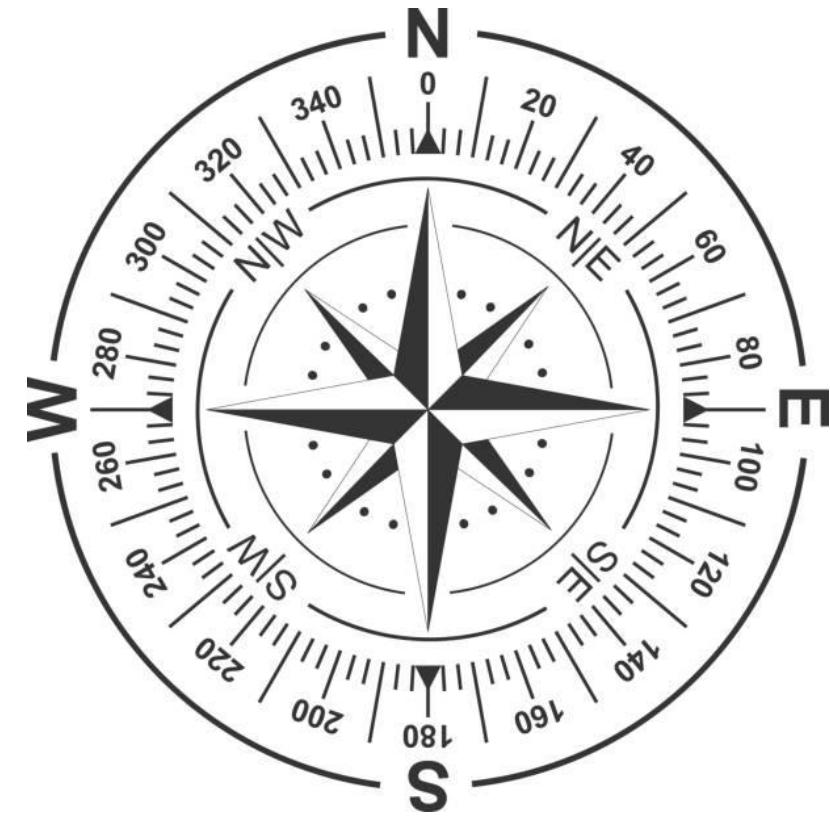
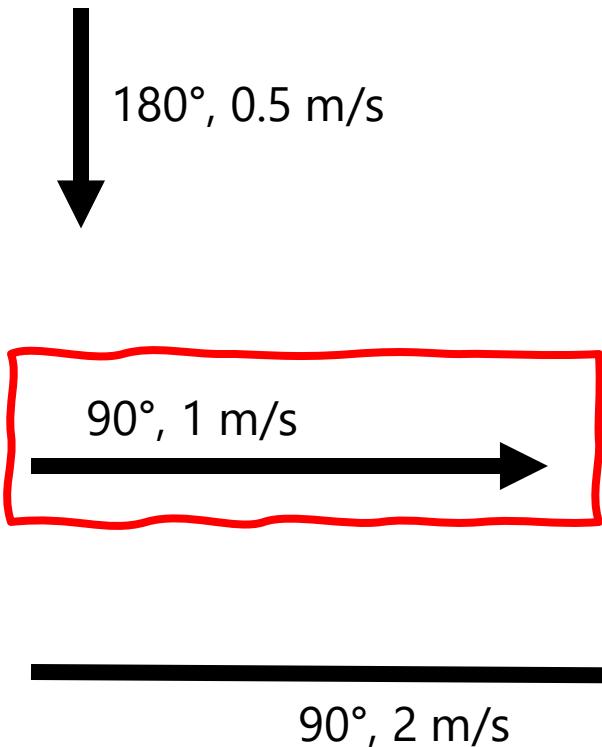
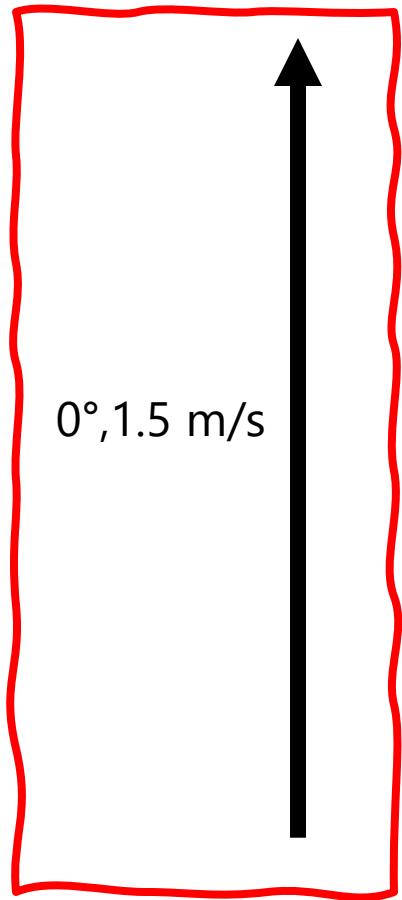
# Vectors



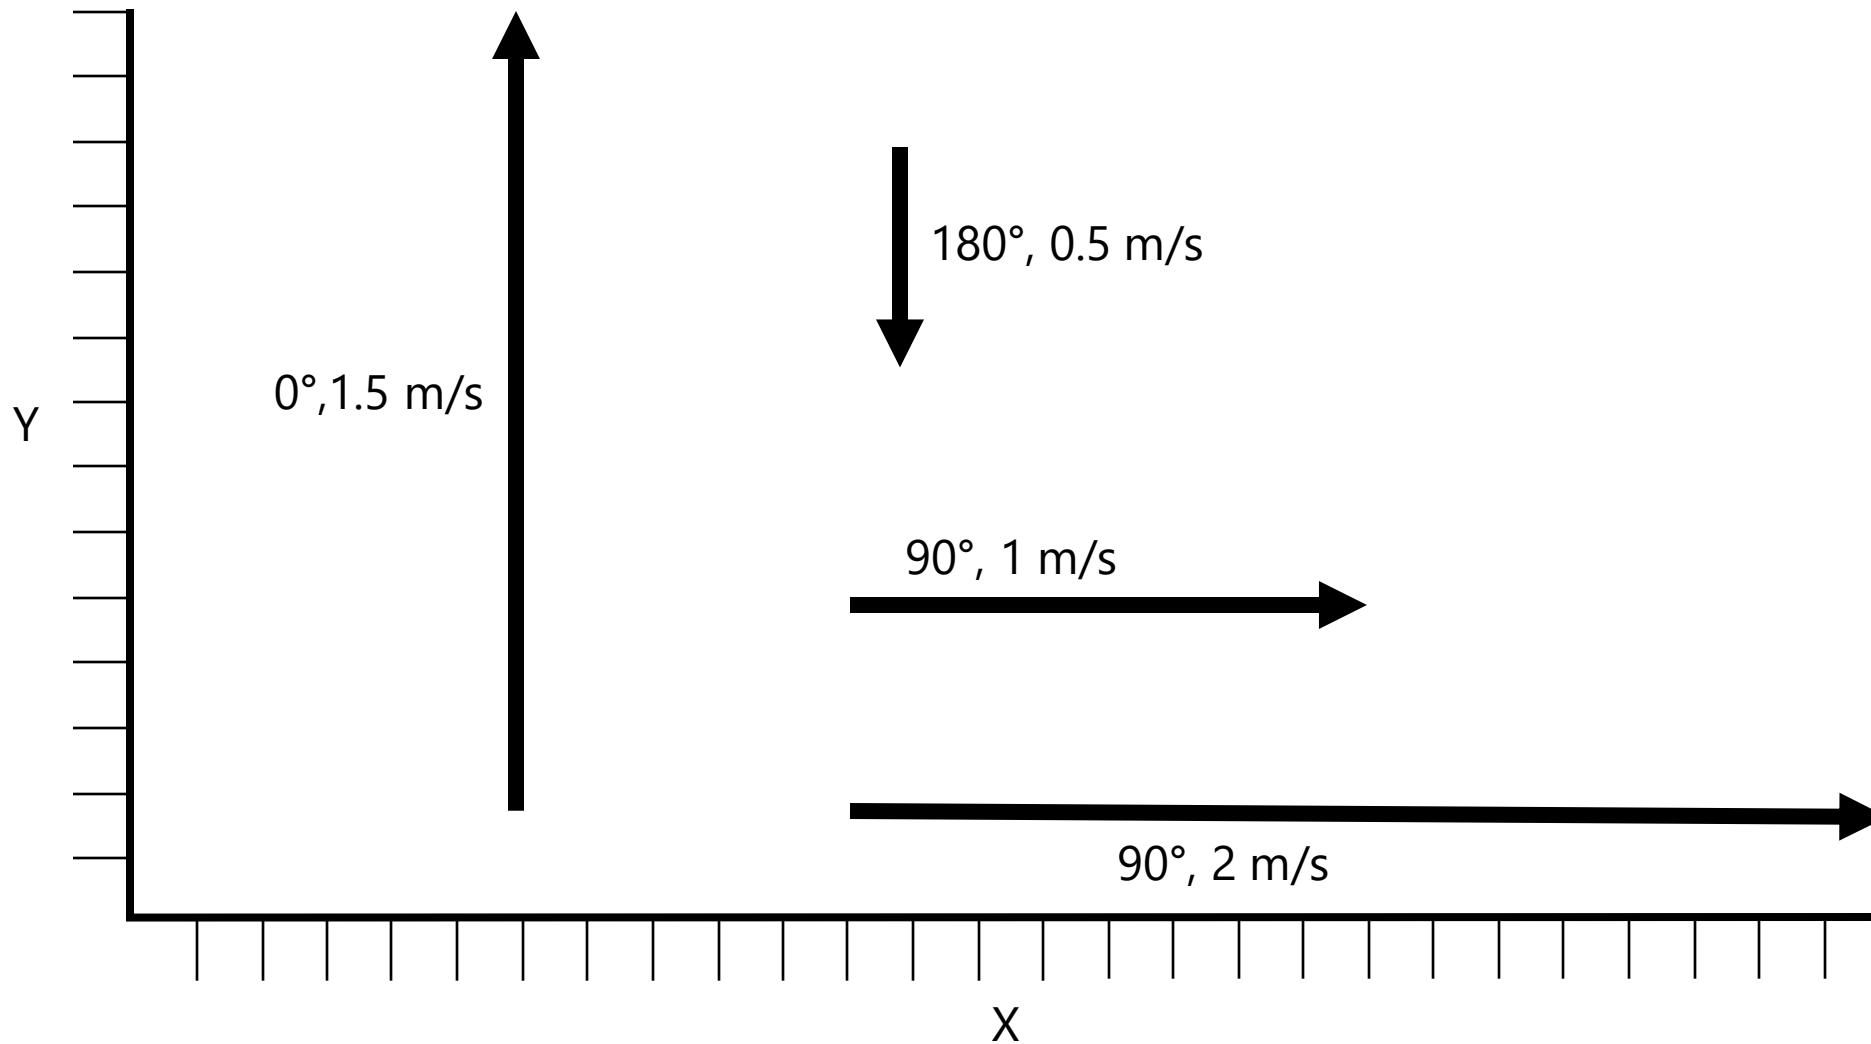
$\downarrow$   
180°, 0.5 m/s  
 $\rightarrow$   
90°, 1 m/s  
 $\rightarrow$   
90°, 2 m/s



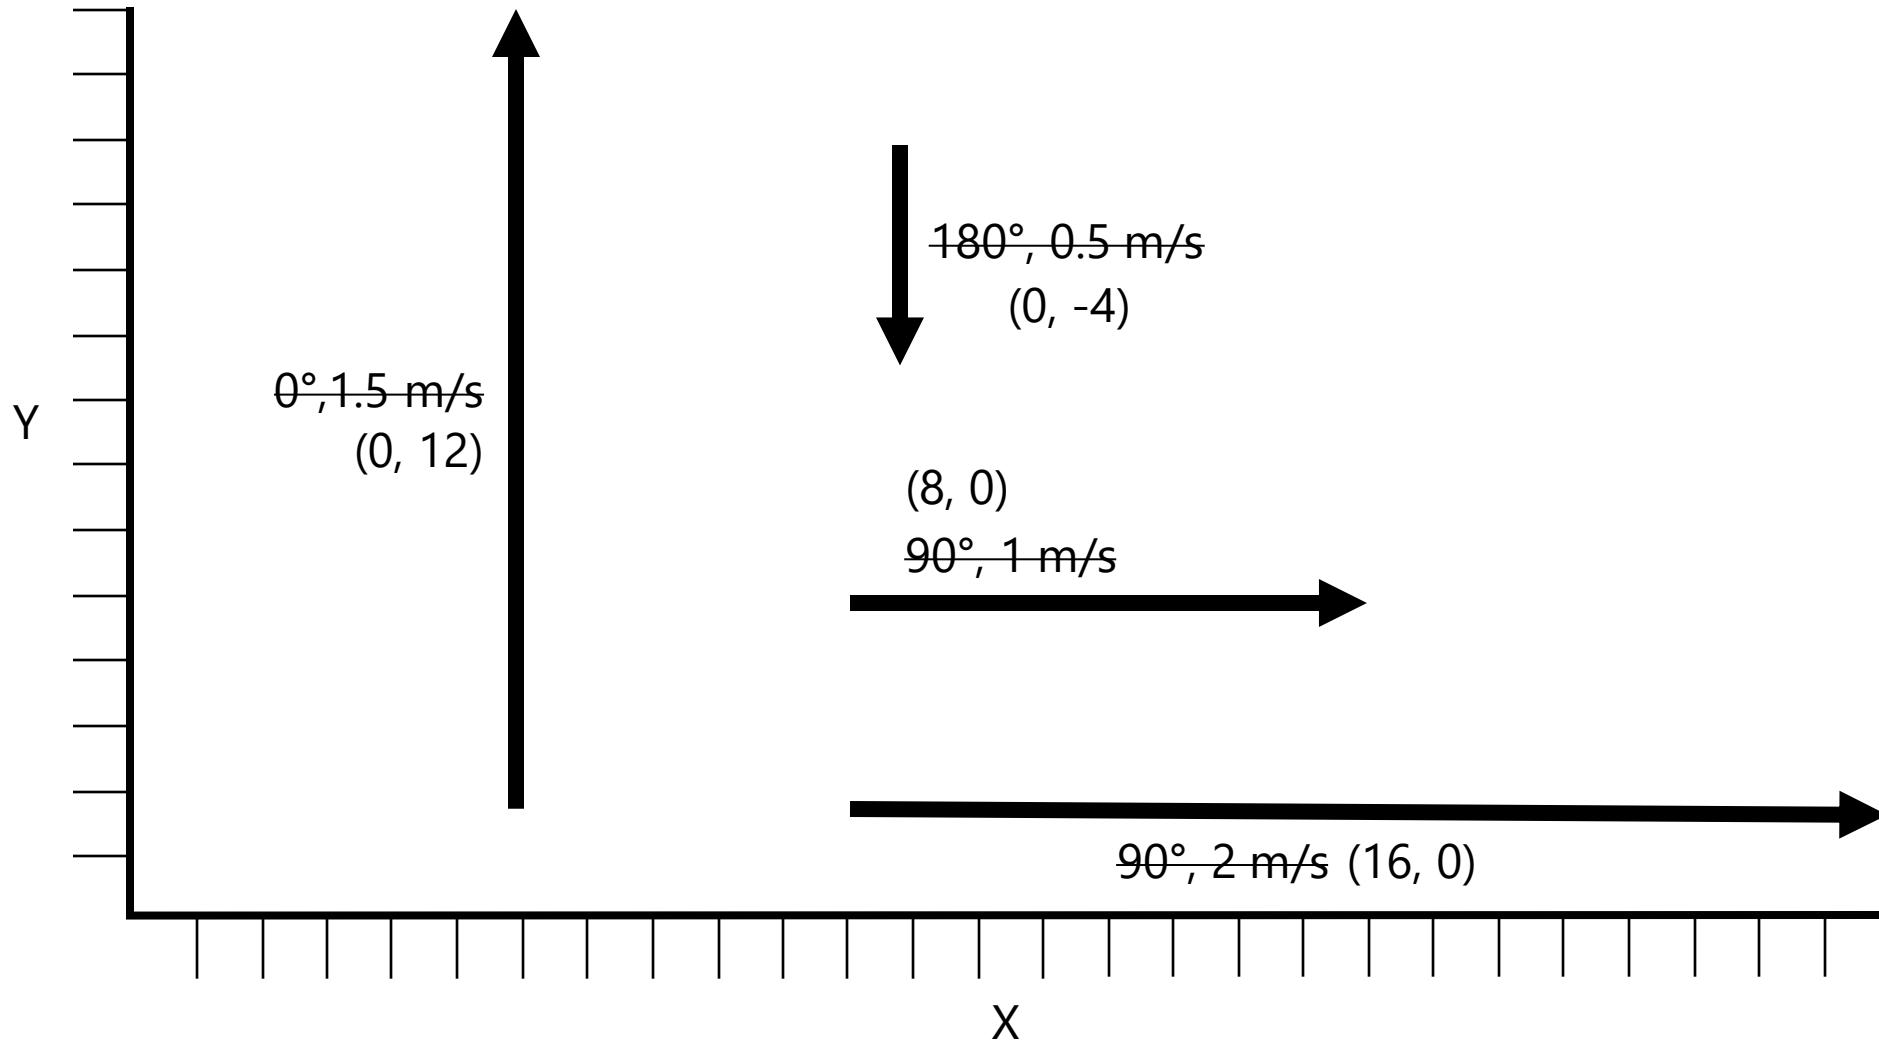
# Vectors



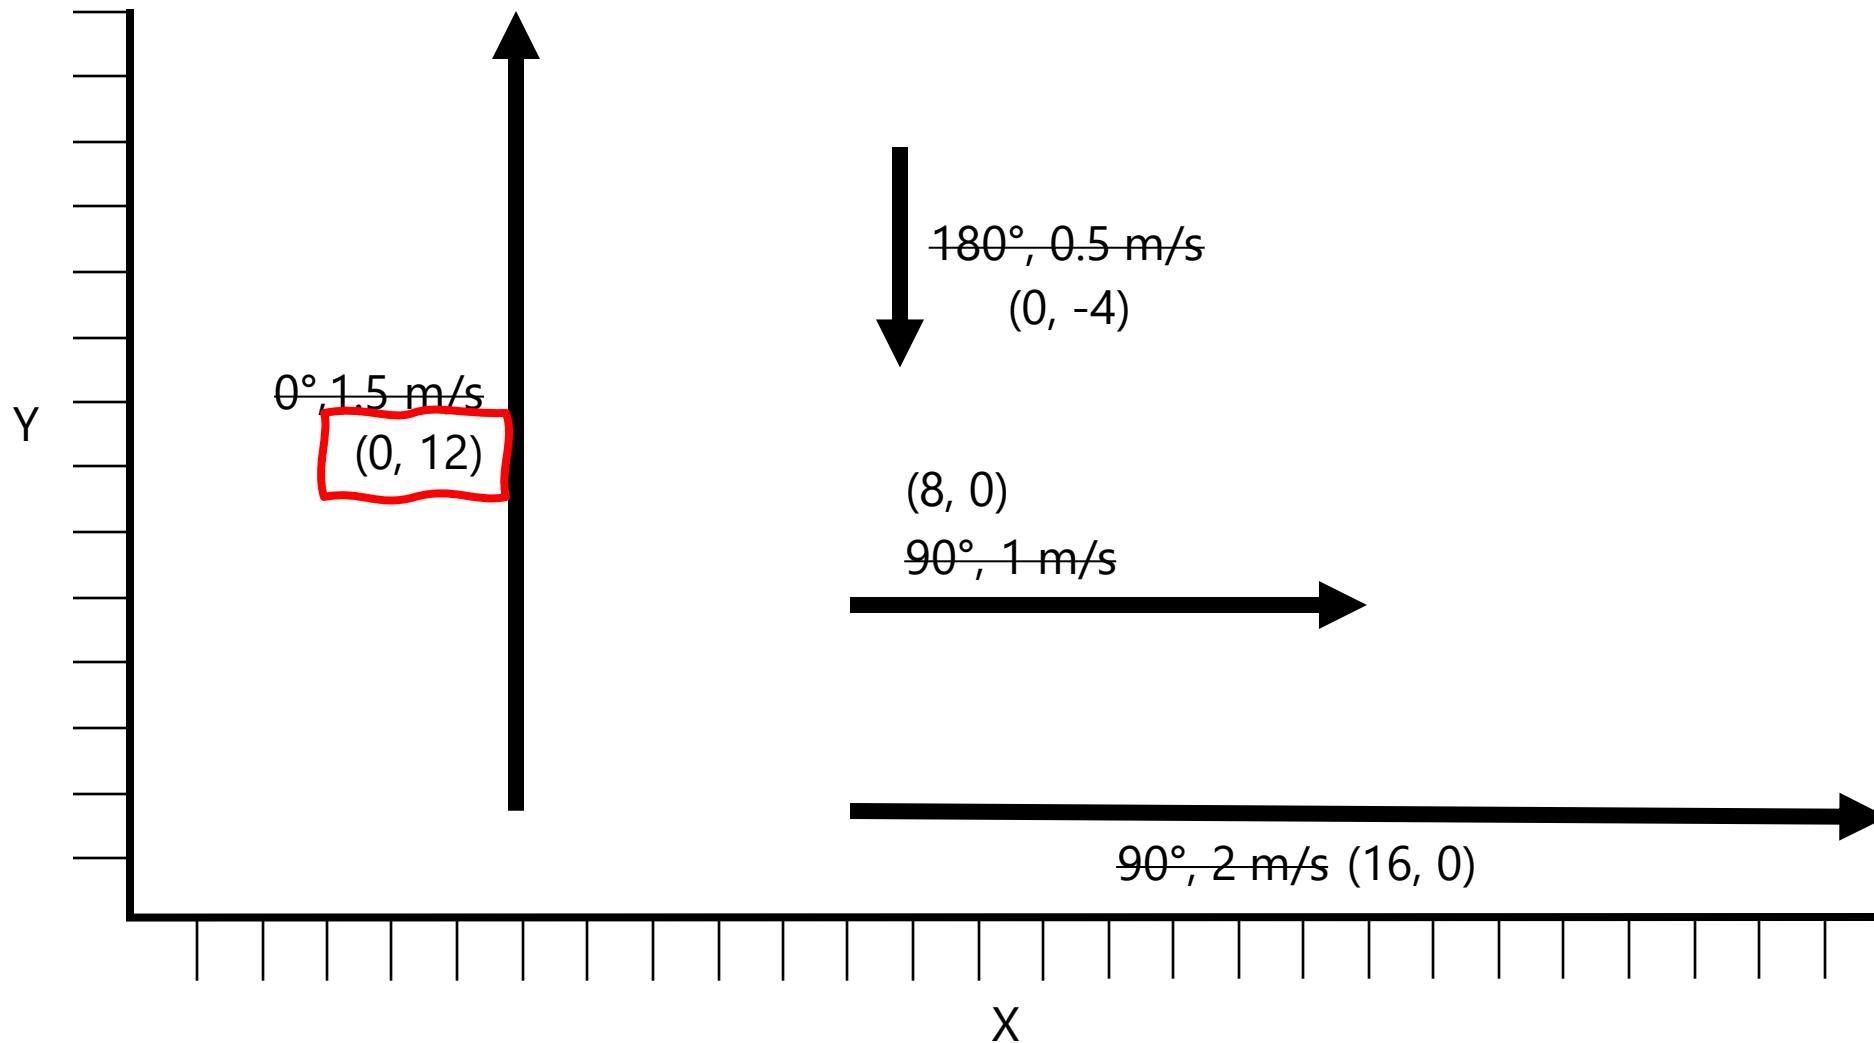
# Vectors



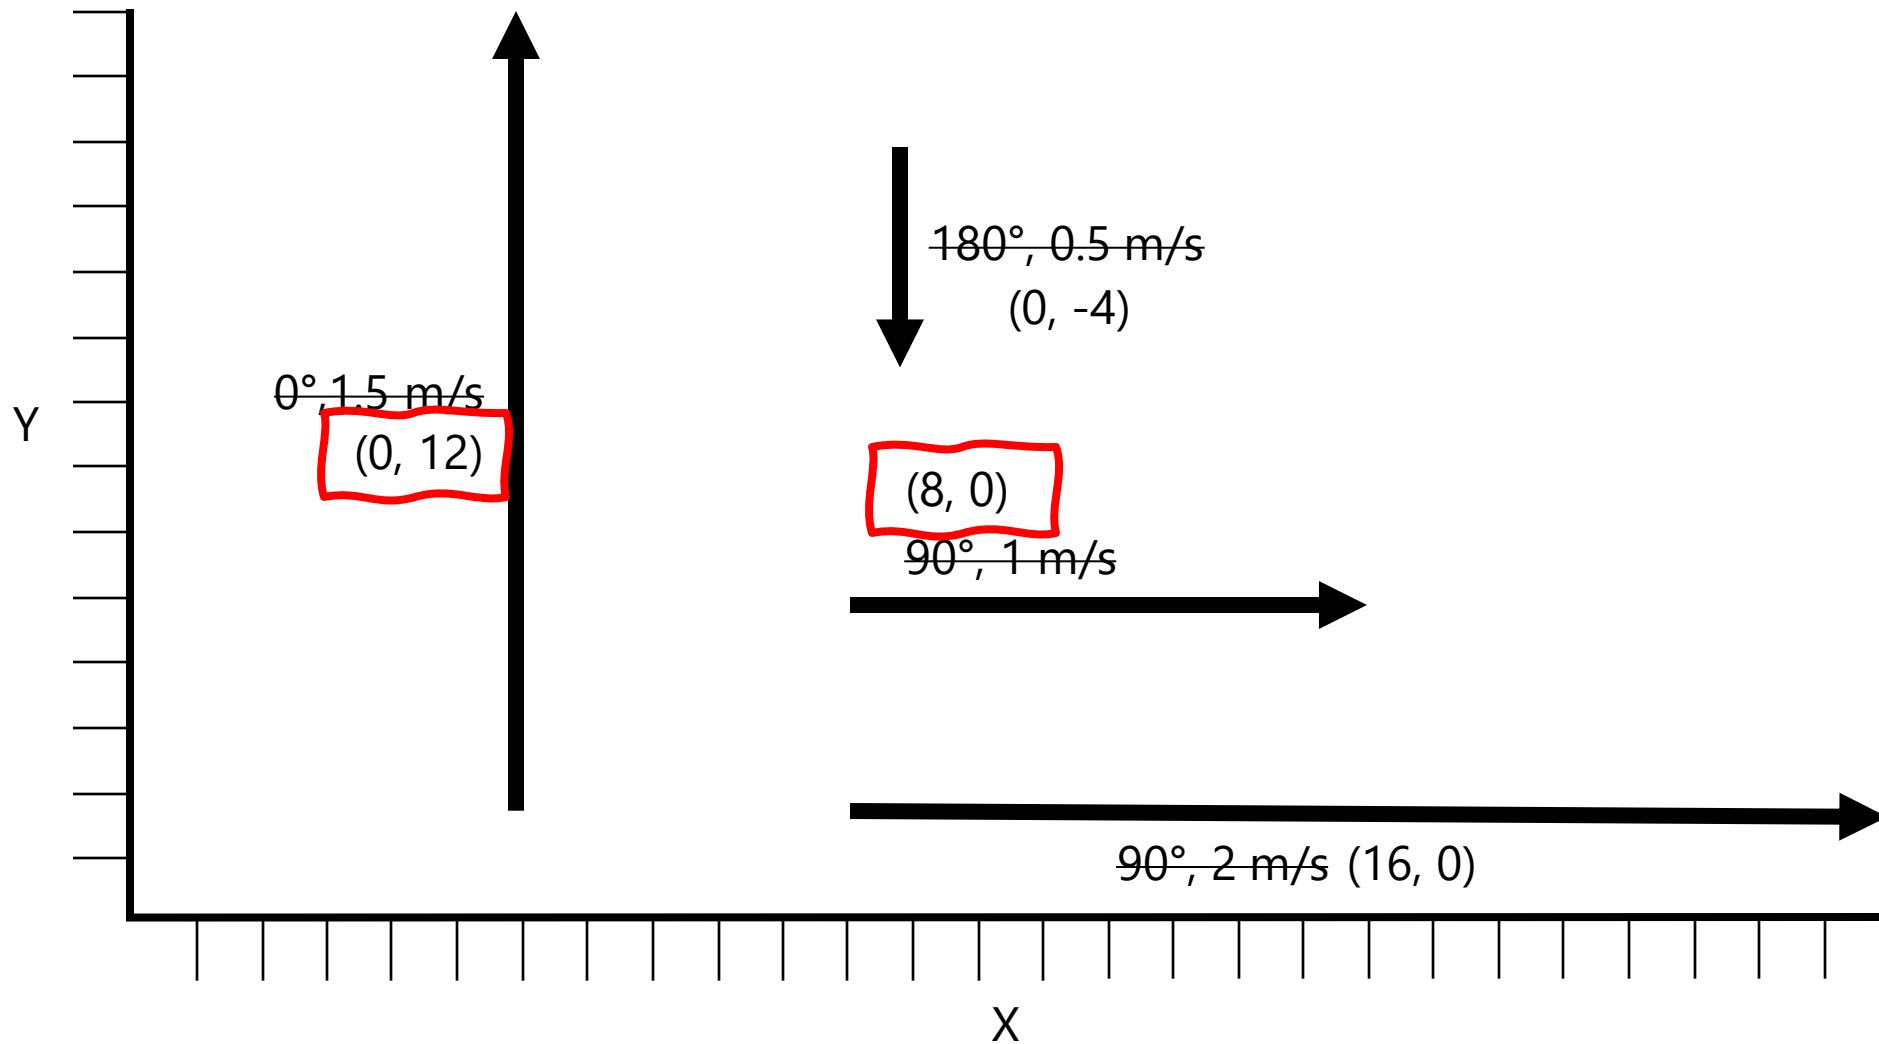
# Vectors



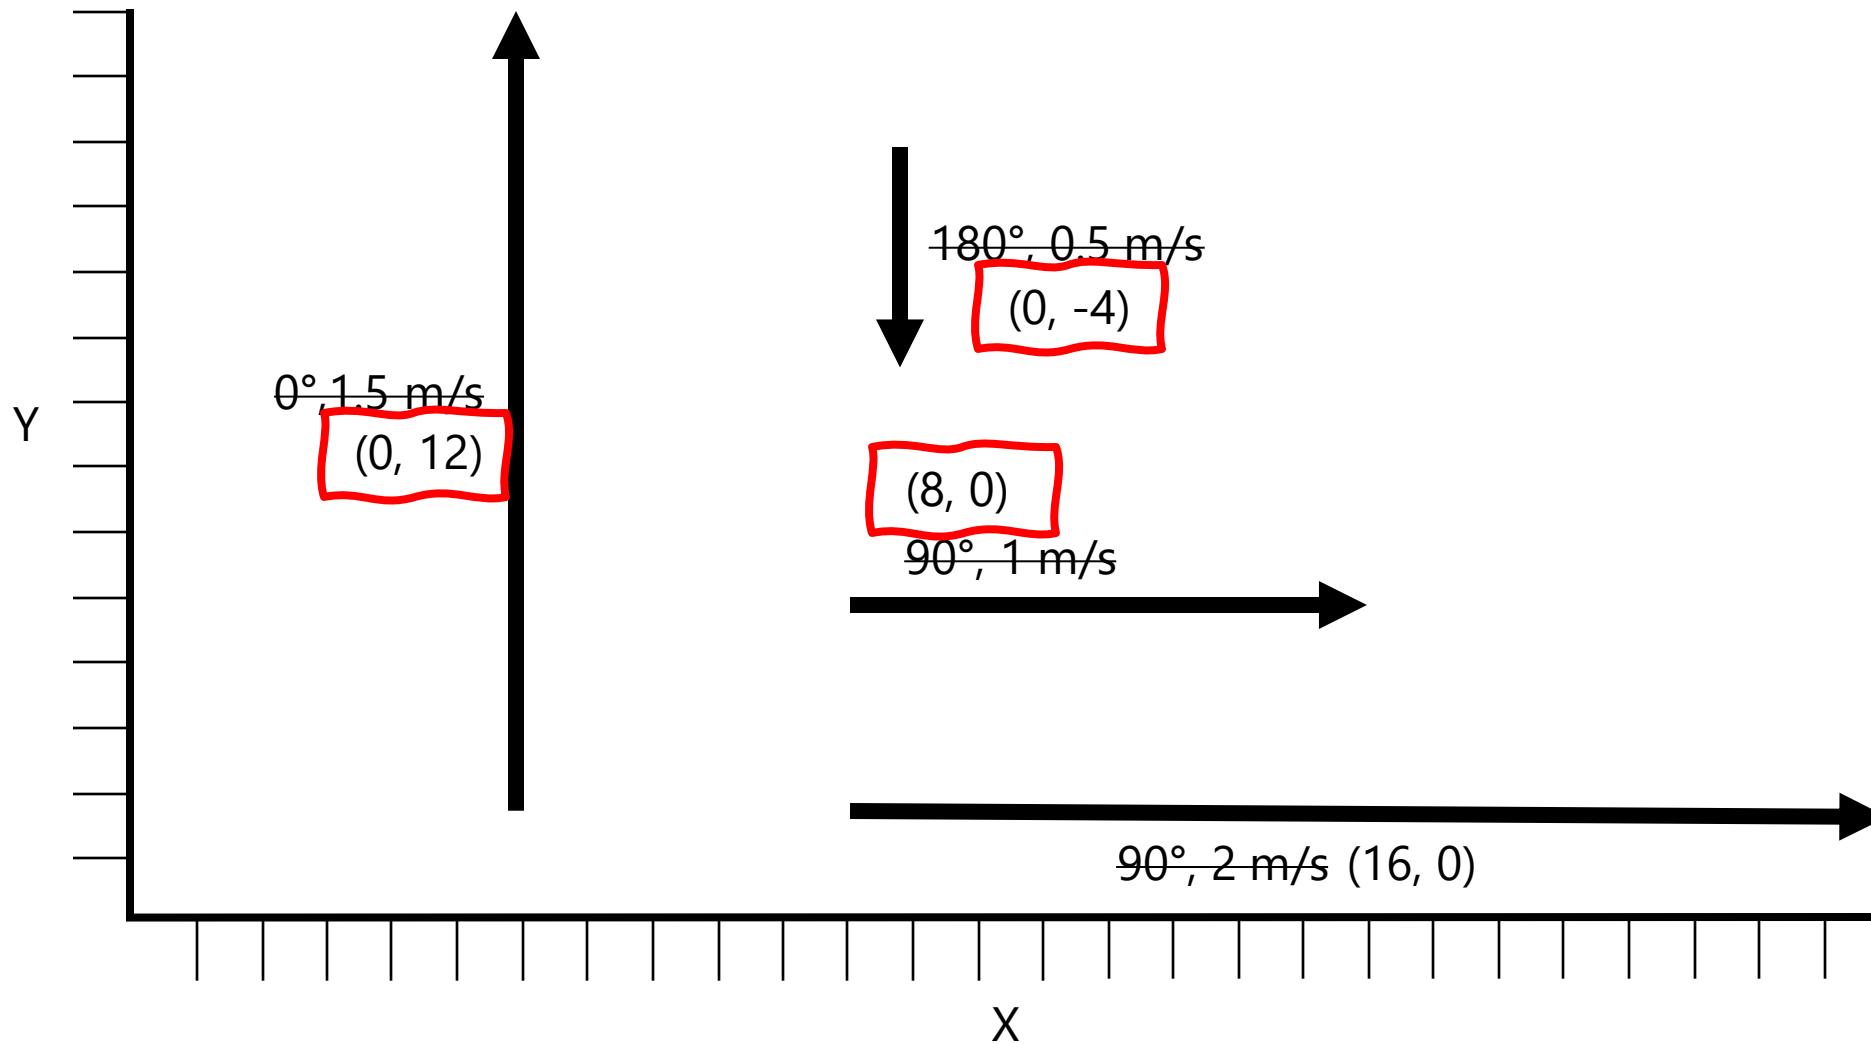
# Vectors



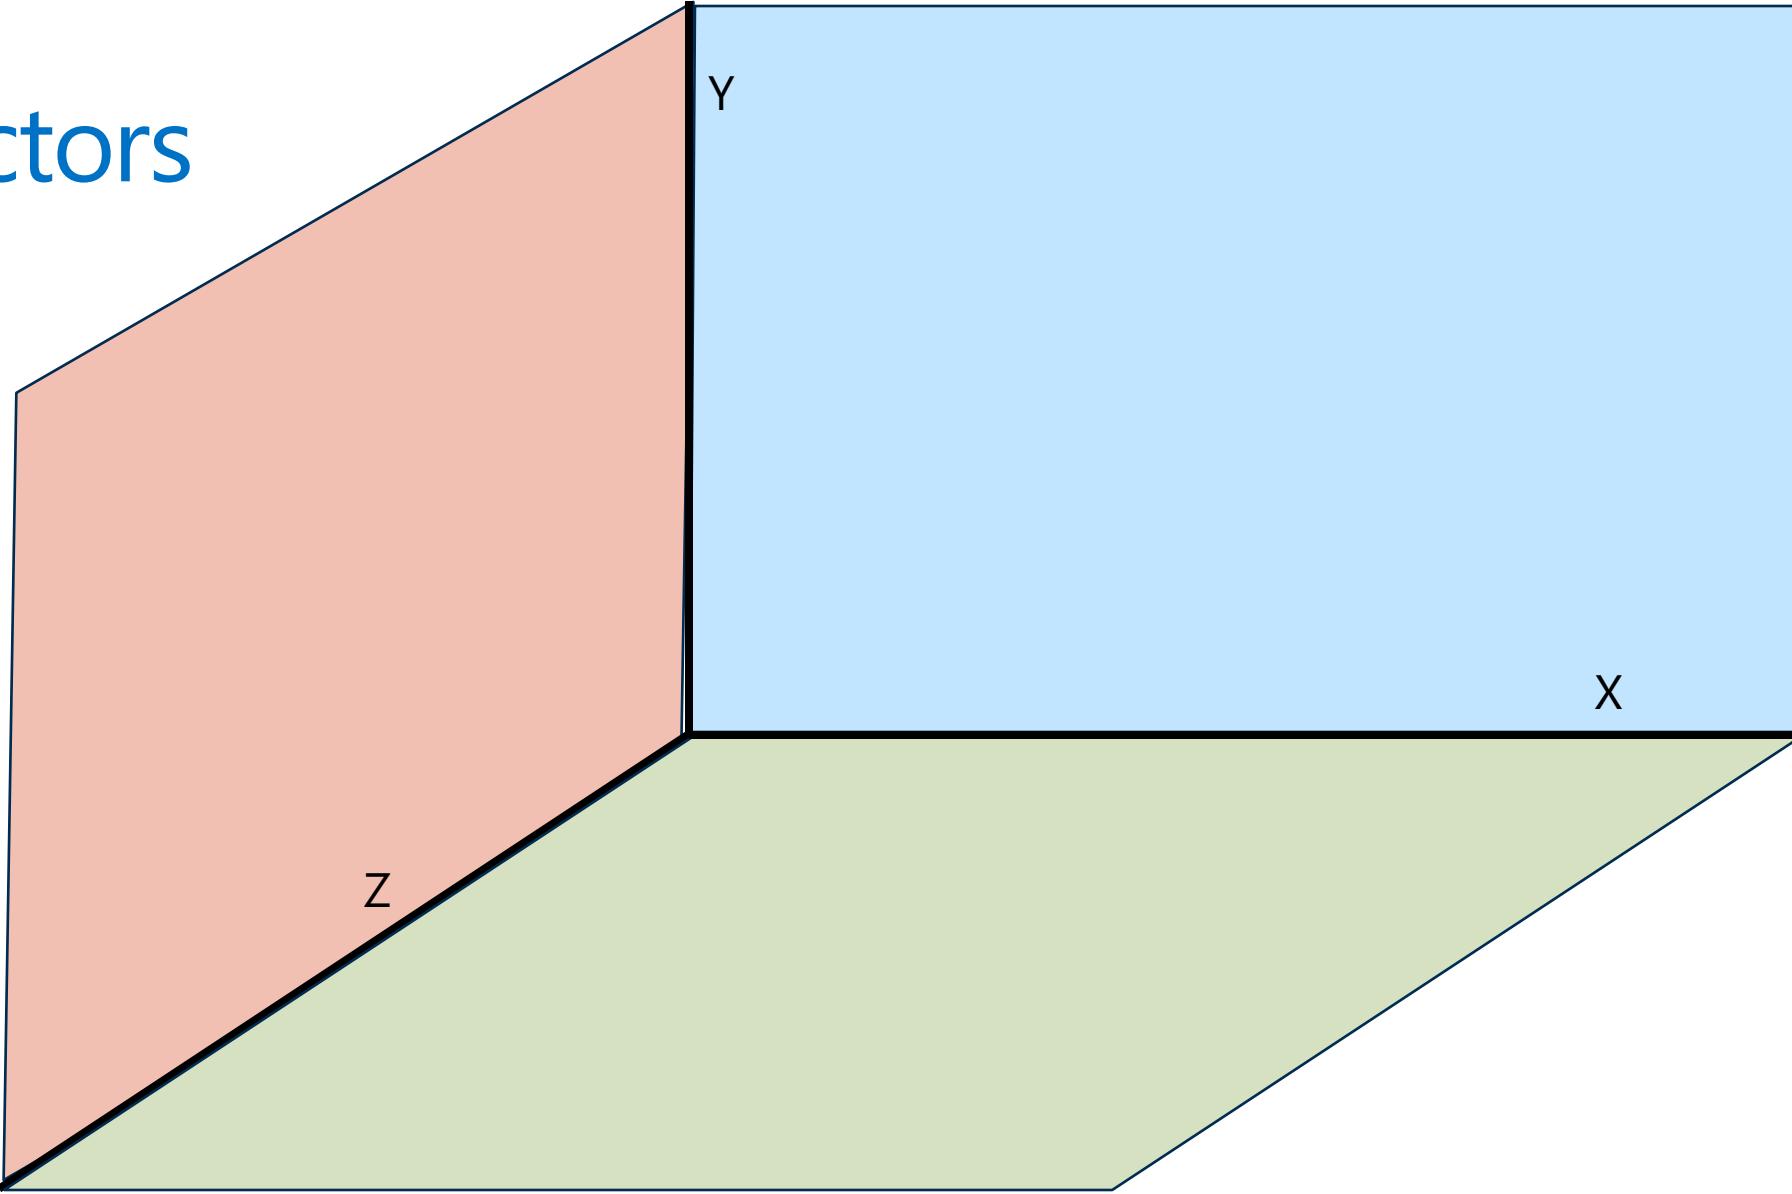
# Vectors



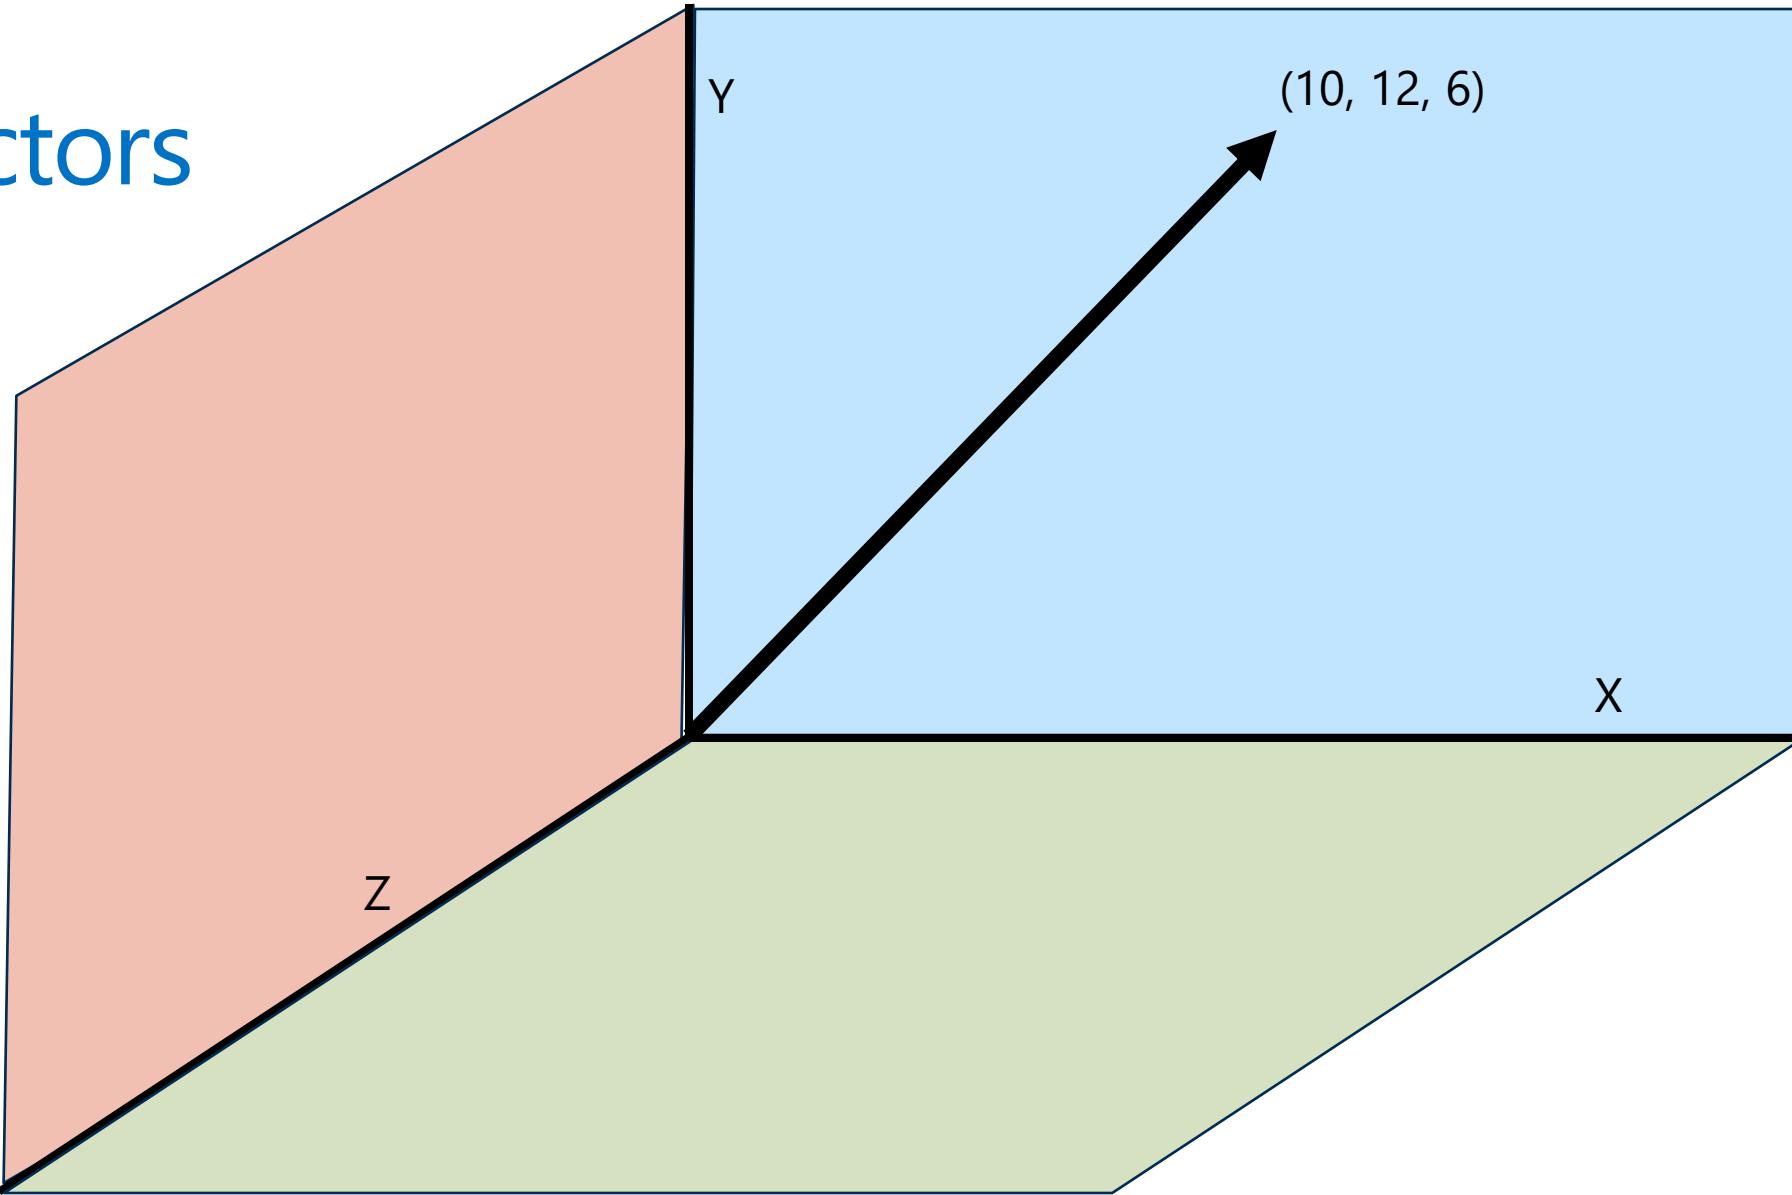
# Vectors



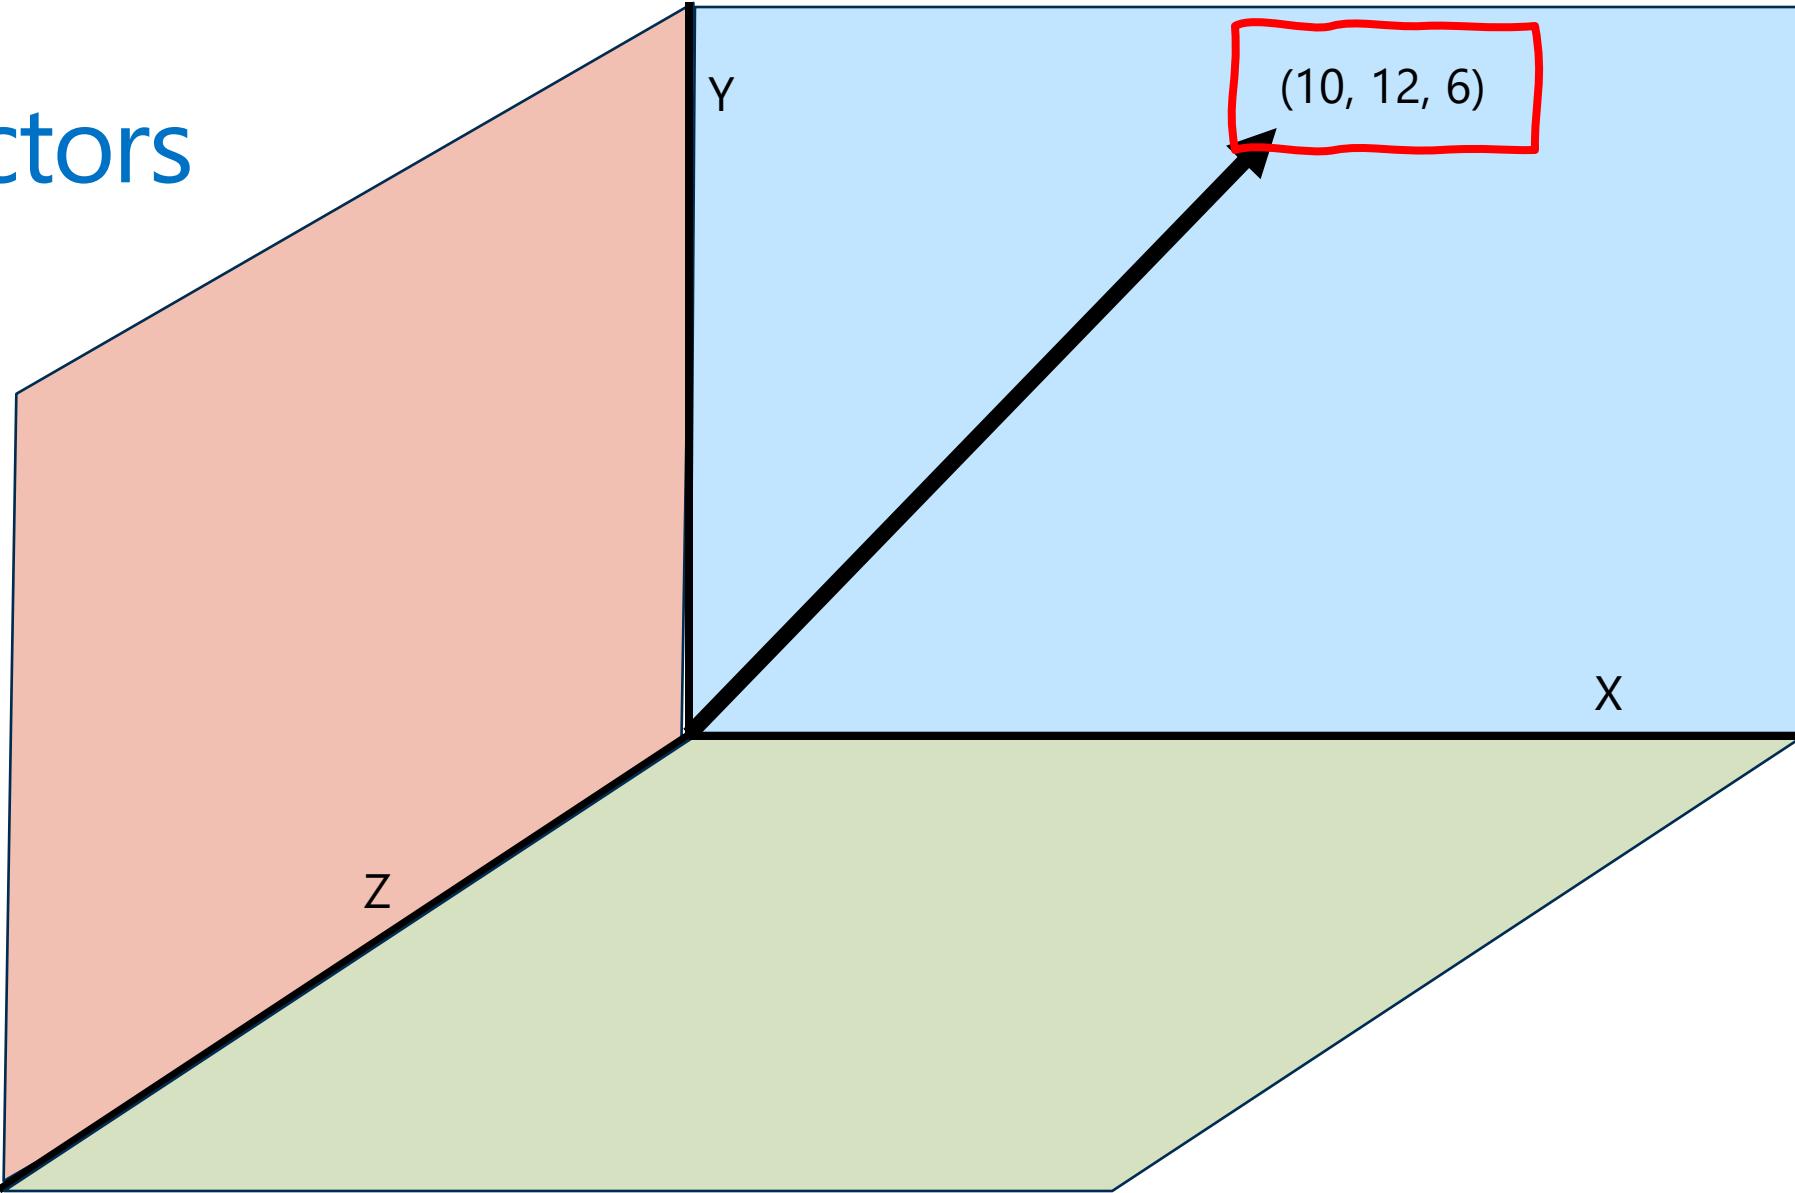
# 3D Vectors



# 3D Vectors



# 3D Vectors



4D+ Vectors

*Sorry*  
**IMAGE**  
NOT AVAILABLE



# 4D+ Vectors

Dimensions	Sample Vector
2	(10, 12)
3	(10, 12, 6)
4	(10, 12, 6, 4)
5	(10, 12, 6, 4, 10)
6	(10, 12, 6, 4, 10, 3)
7	(10, 12, 6, 4, 10, 3, 144)
N	(10, 12, 6, 4, 10, 3, 144, ...)

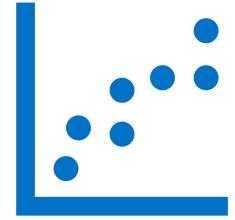
# 4D+ Vectors

Dimensions	Sample Vector
2	(10, 12)
3	(10, 12, 6)
4	(10, 12, 6, 4)
5	(10, 12, 6, 4, 10)
6	(10, 12, 6, 4, 10, 3)
7	(10, 12, 6, 4, 10, 3, 144)
N	(10, 12, 6, 4, 10, 3, 144, ...)

**LLMS today:  
384 to 3,000  
dimension vectors**

# Embeddings

# Embeddings

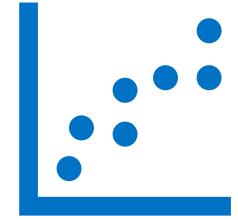


Storing meaning using vectors

An **embedding** is just a **vector** that points in a direction **representing meaning**

The **closer** two embeddings point in the **same direction**, the more **similar** their meaning

# Embeddings



Storing meaning using vectors

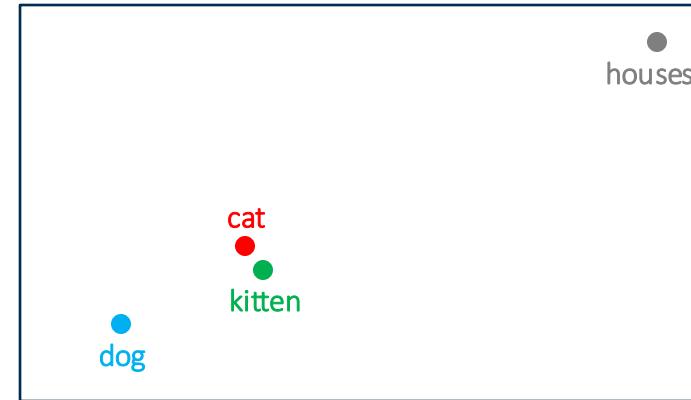
An **embedding** is just a **vector** that points in a direction **representing meaning**

The **closer** two embeddings point in the **same direction**, the more **similar** their meaning

# Embeddings

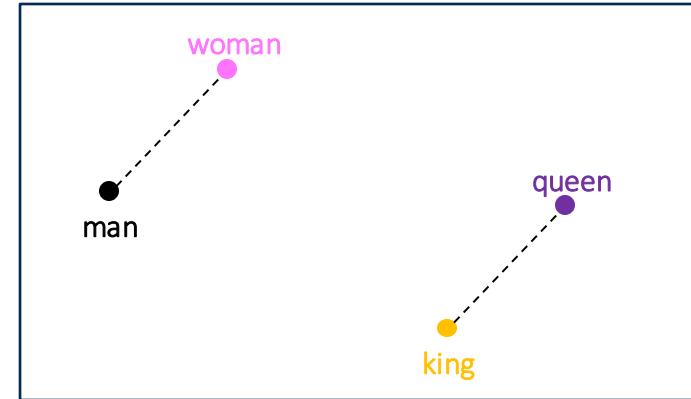
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization



<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

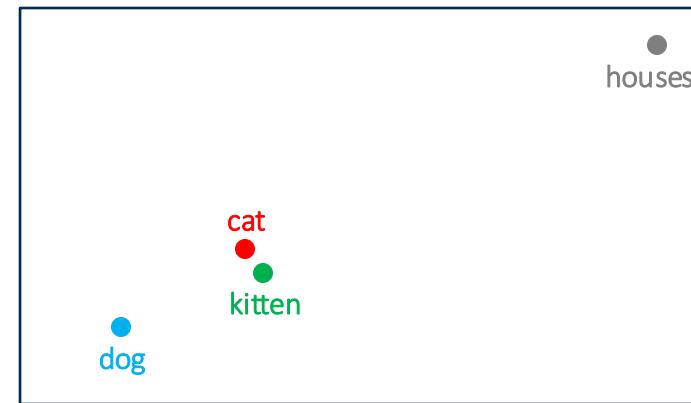
Simplified for  
Visualization



# Embeddings

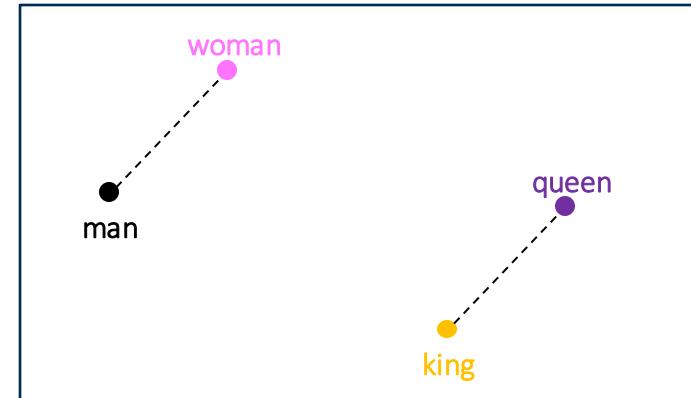
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization



	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

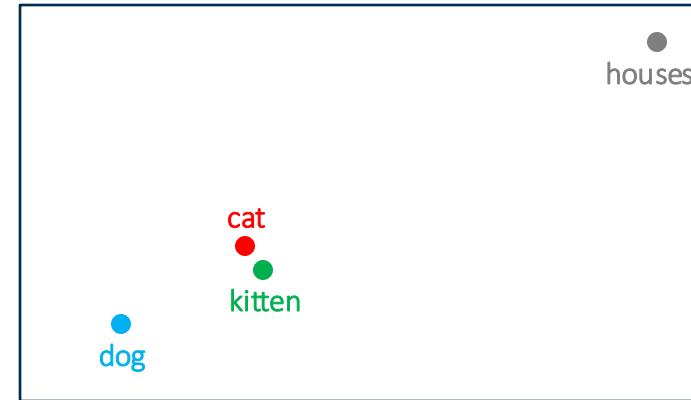
Simplified for  
Visualization



# Embeddings

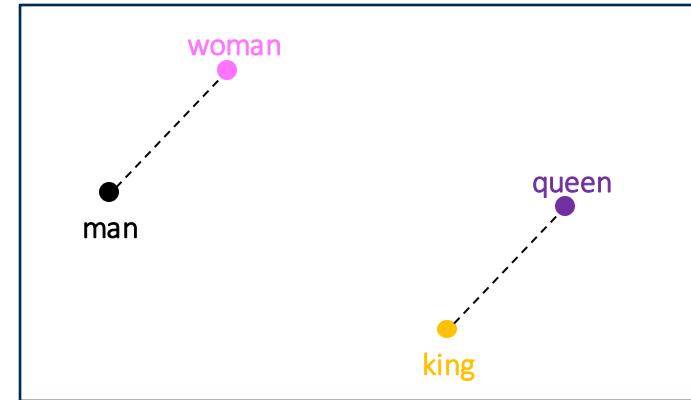
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization

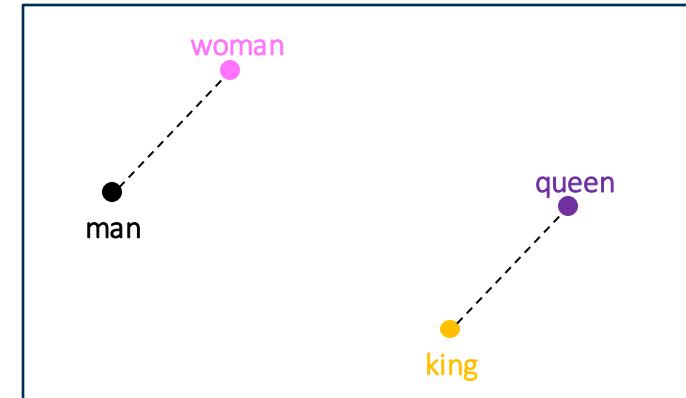
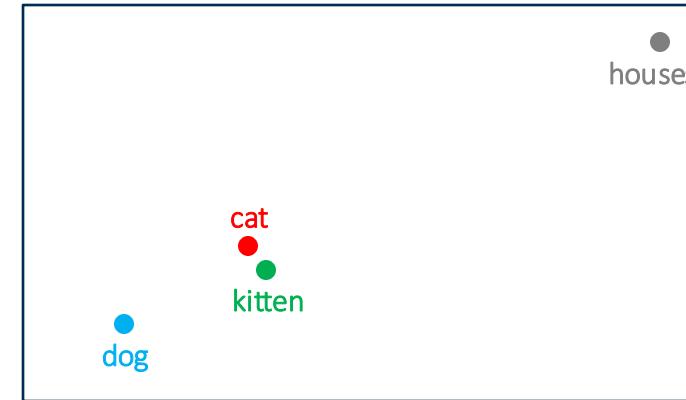
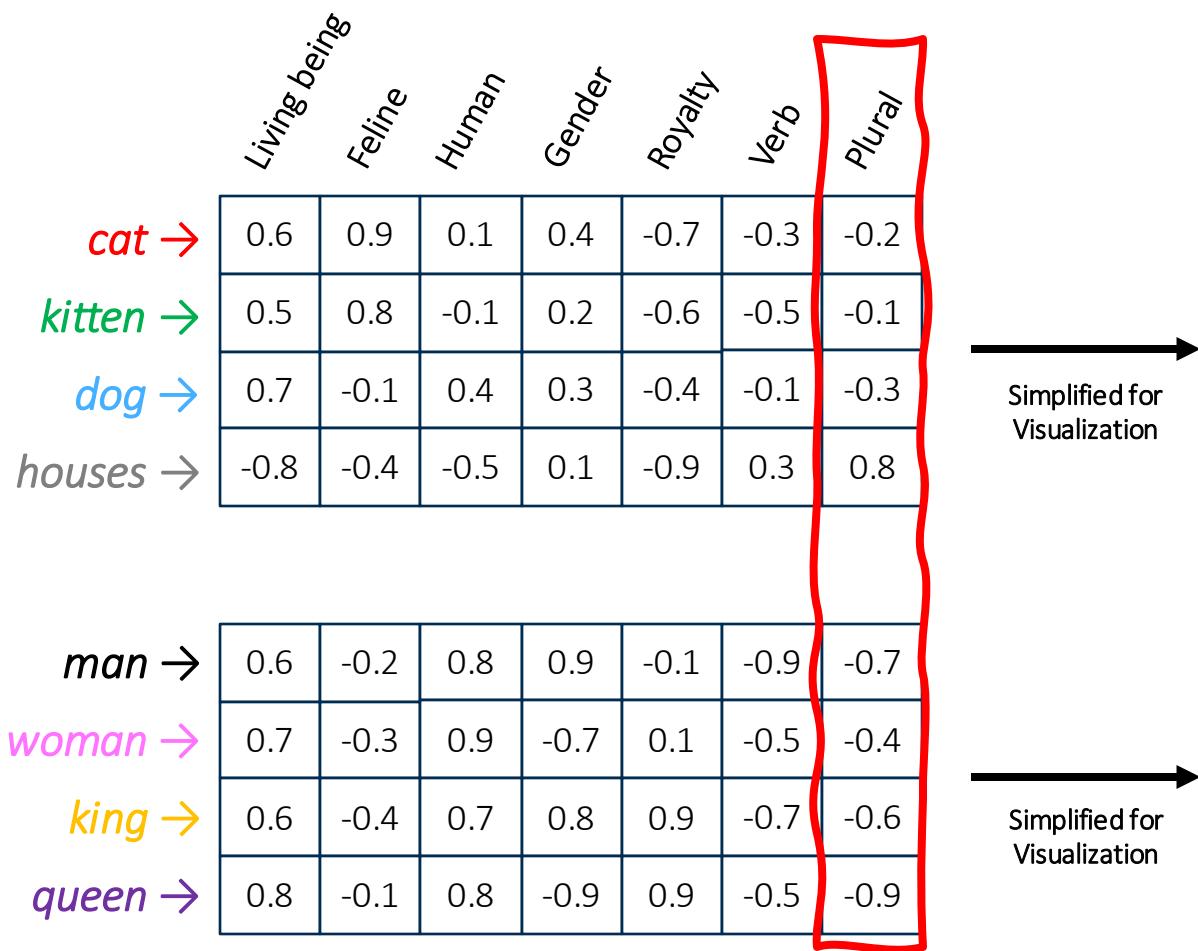


	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

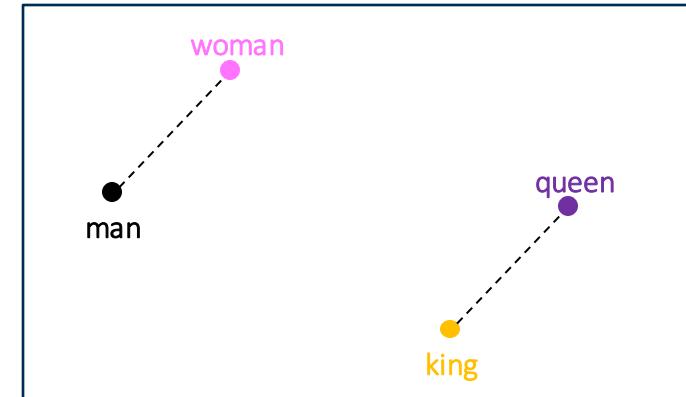
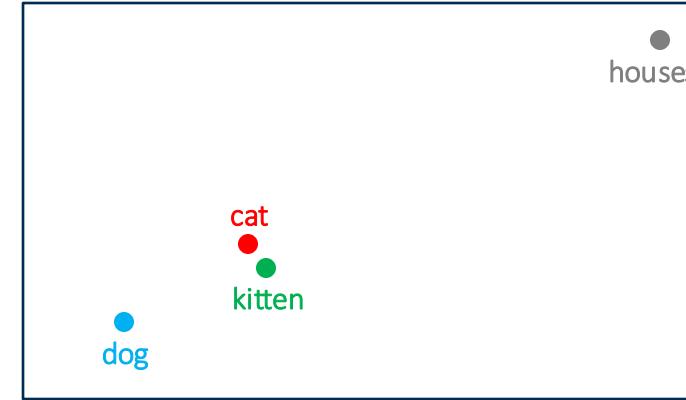
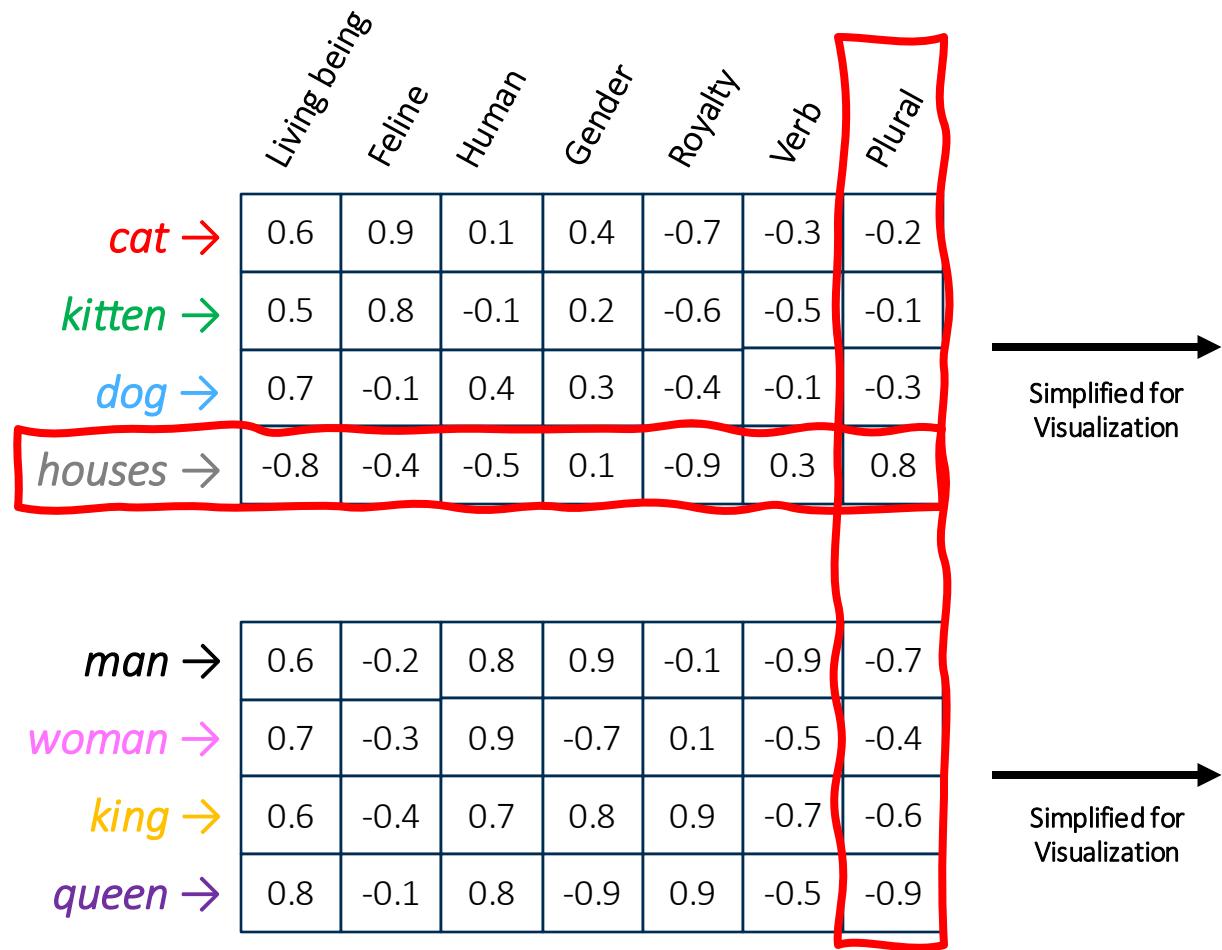
Simplified for  
Visualization



# Embeddings



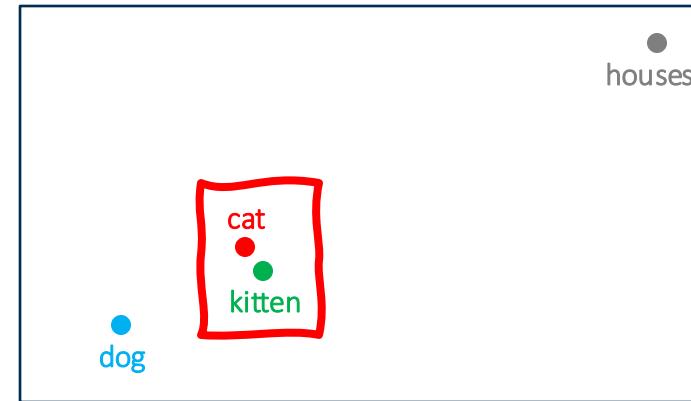
# Embeddings



# Embeddings

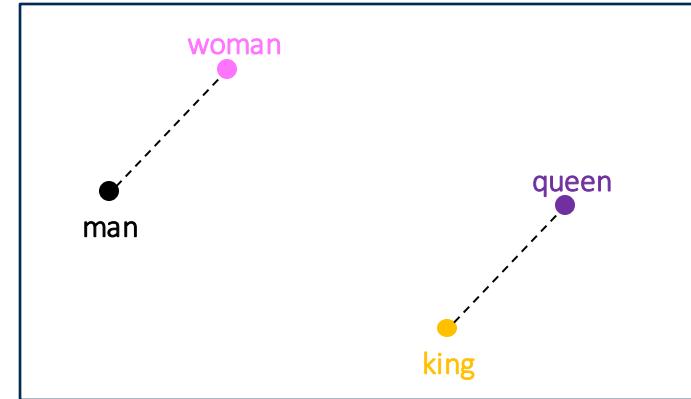
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization



<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

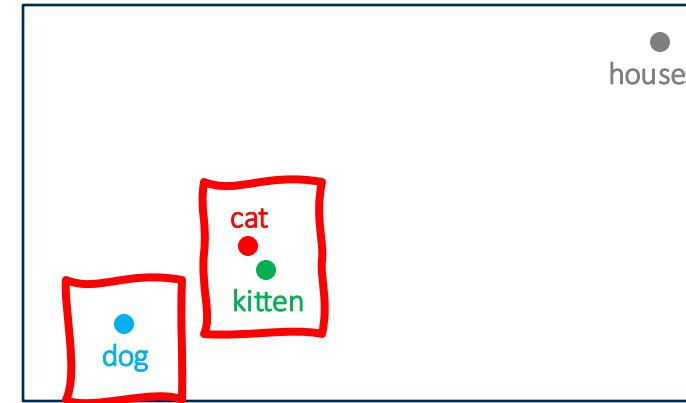
Simplified for  
Visualization



# Embeddings

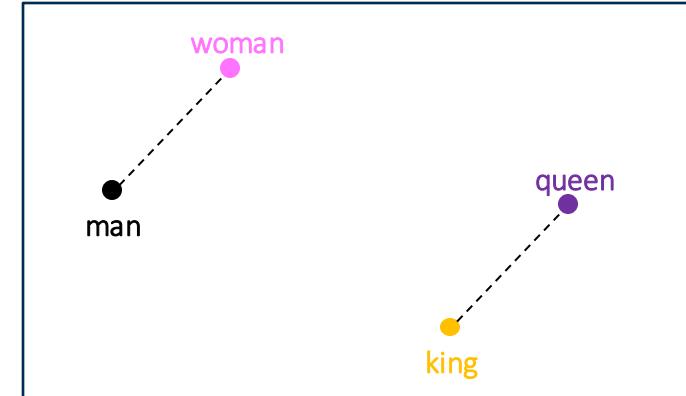
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization

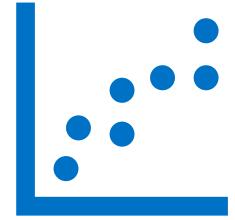


<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

Simplified for  
Visualization



# Embeddings



Storing meaning using vectors

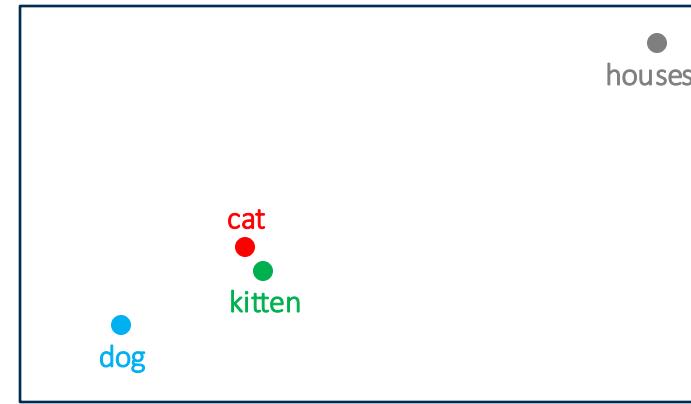
An **embedding** is just a **vector** that points in a direction **representing meaning**

The **closer** two embeddings point in the **same direction**, the more **similar their meaning**

# Embeddings

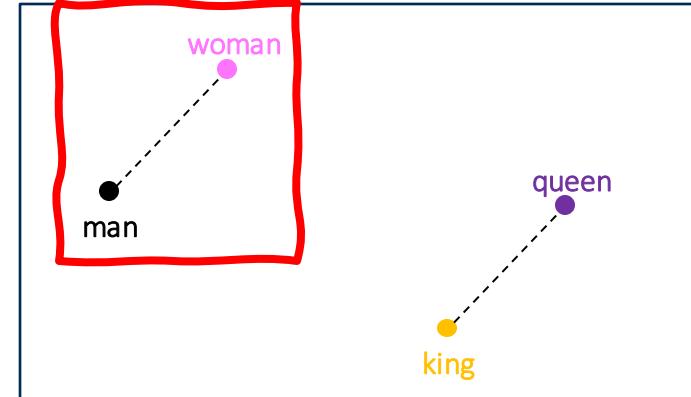
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization



<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

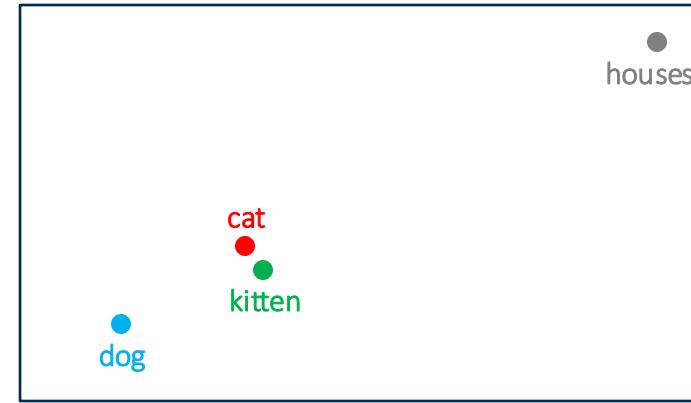
Simplified for  
Visualization



# Embeddings

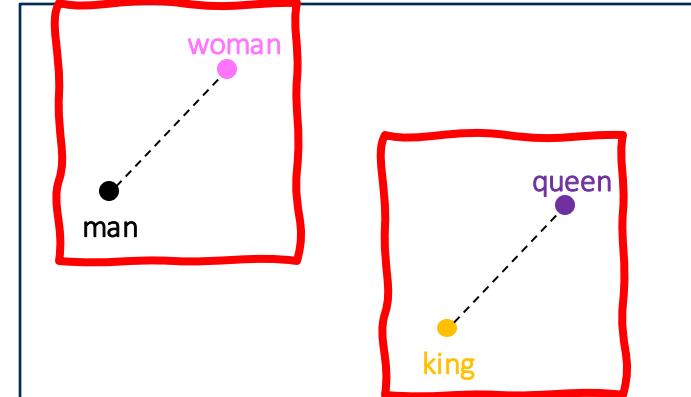
	Living being	Feline	Human	Gender	Royalty	Verb	Plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

Simplified for  
Visualization



<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	-0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.6	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.9	-0.5	-0.9

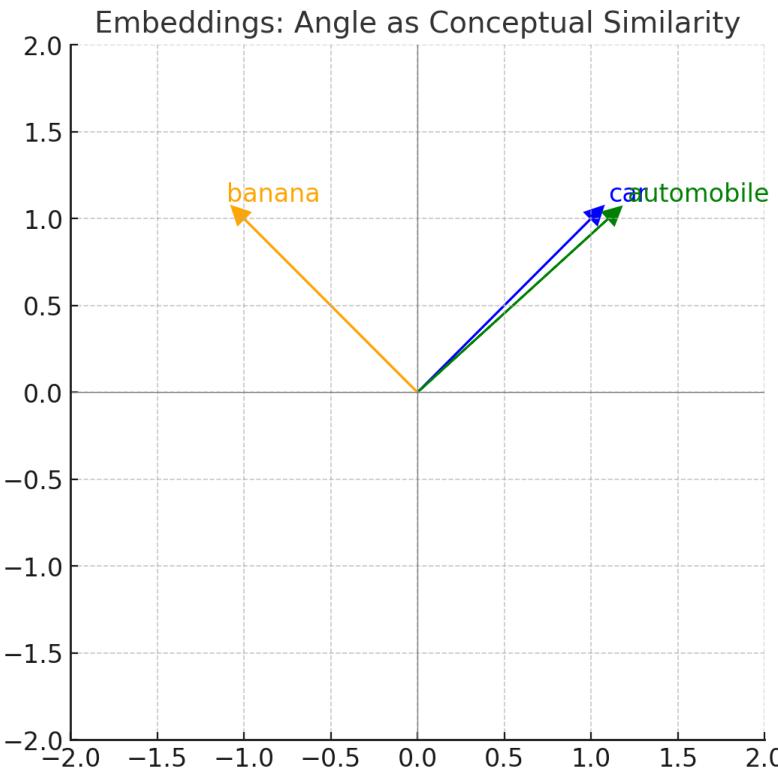
Simplified for  
Visualization



# Cosine Similarity

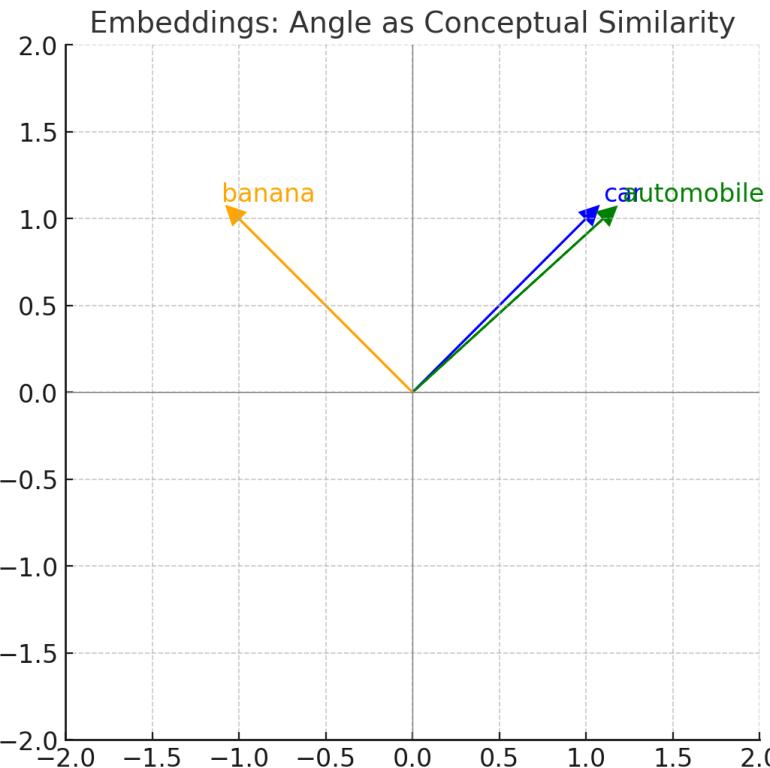
# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

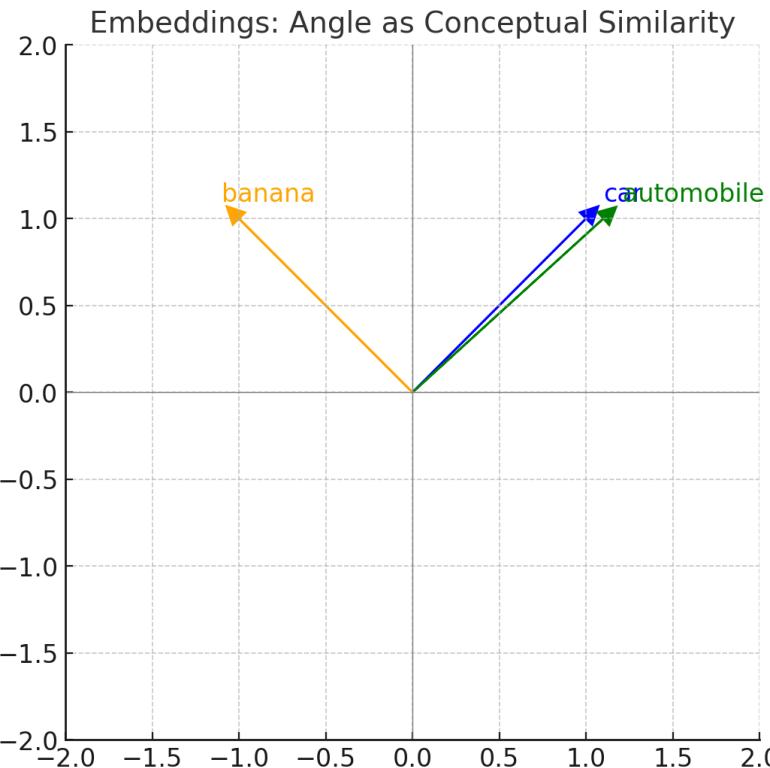
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

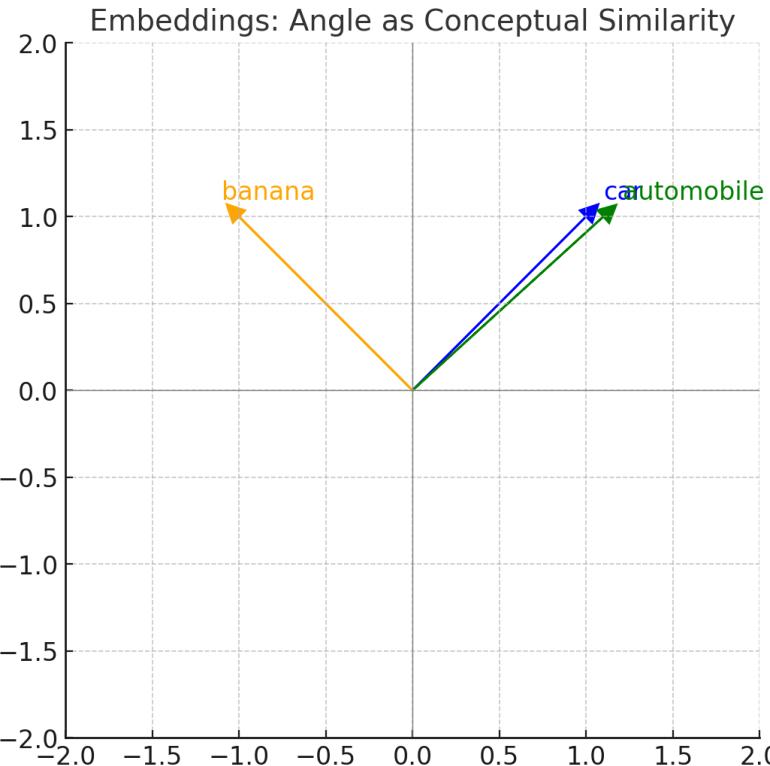
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

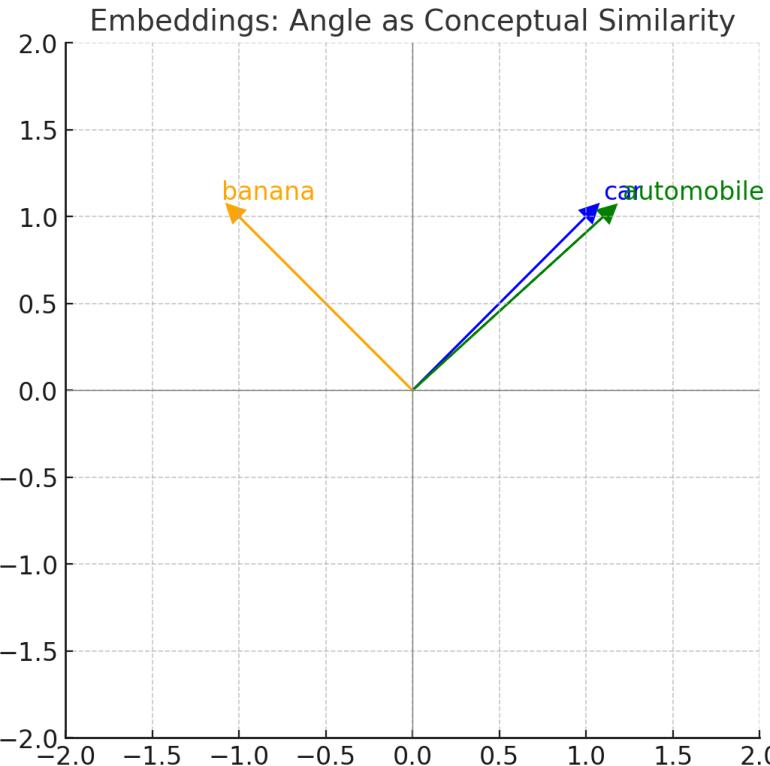
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

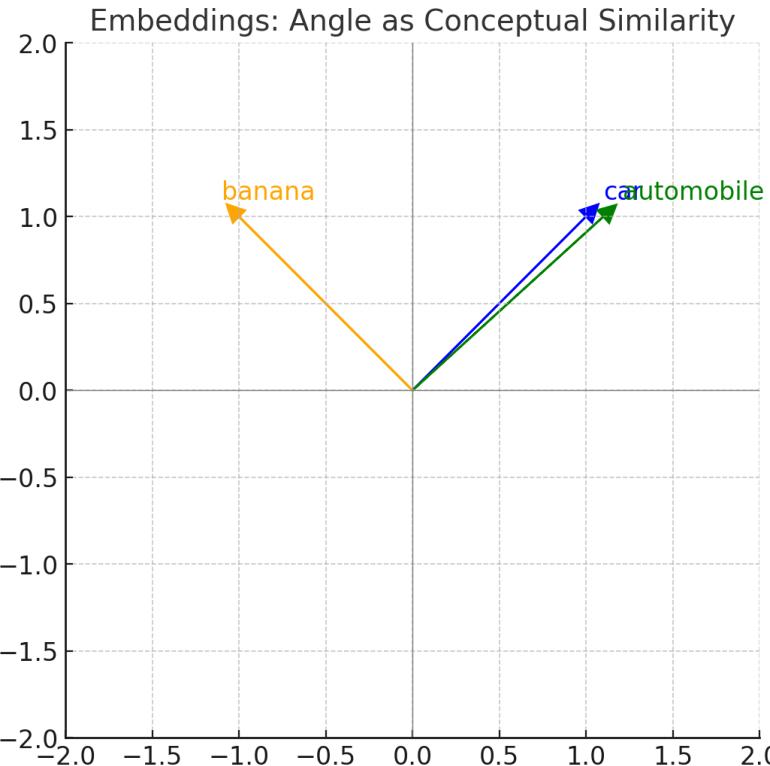
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

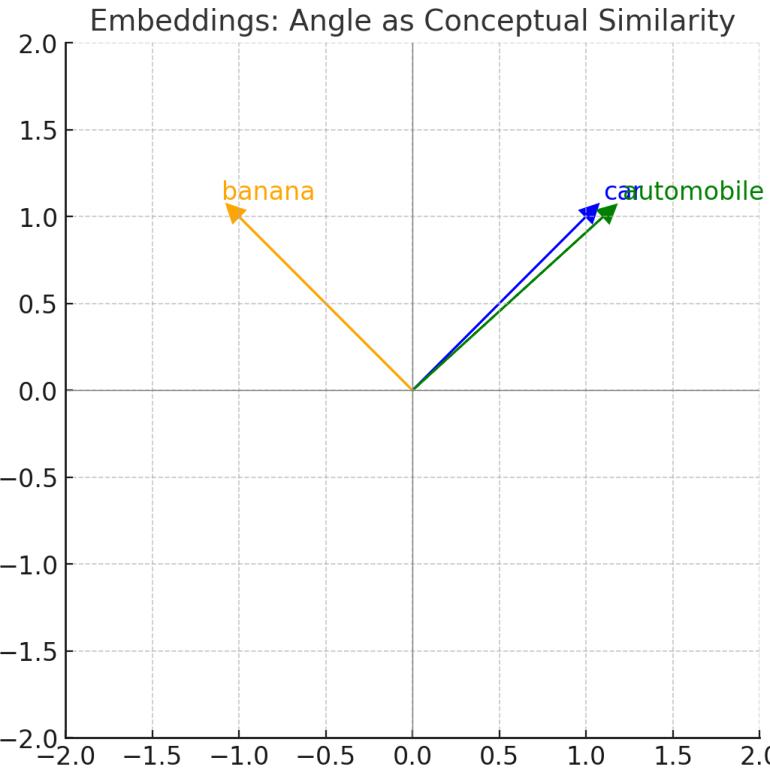
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

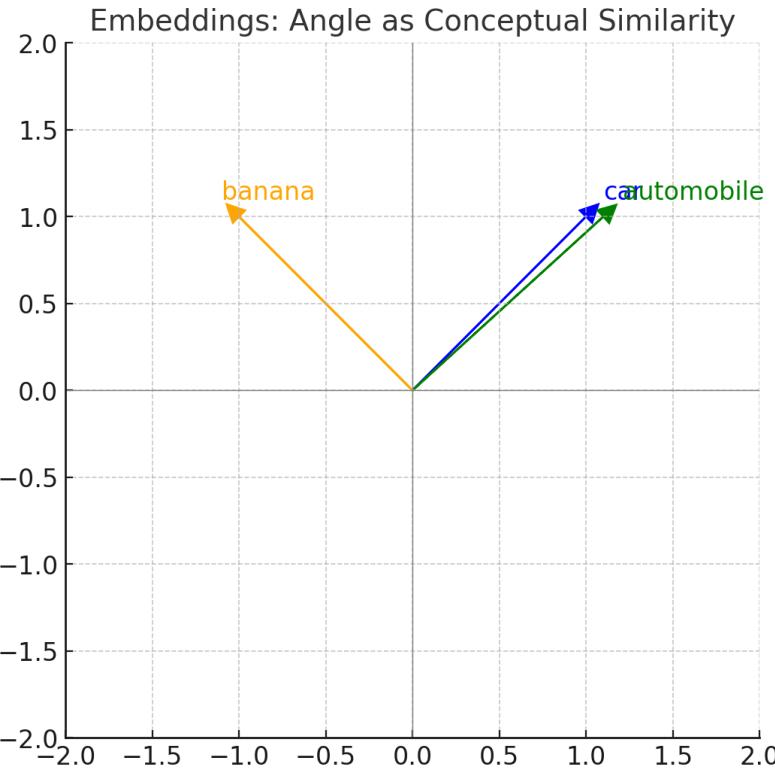
$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$

# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

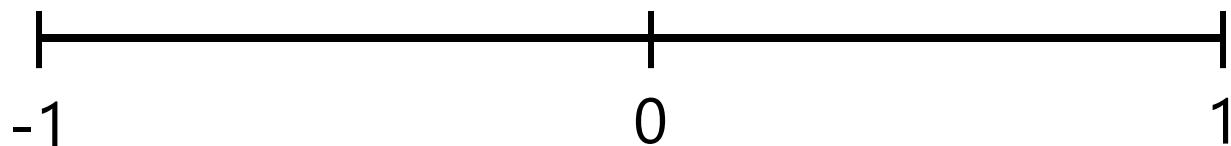
$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

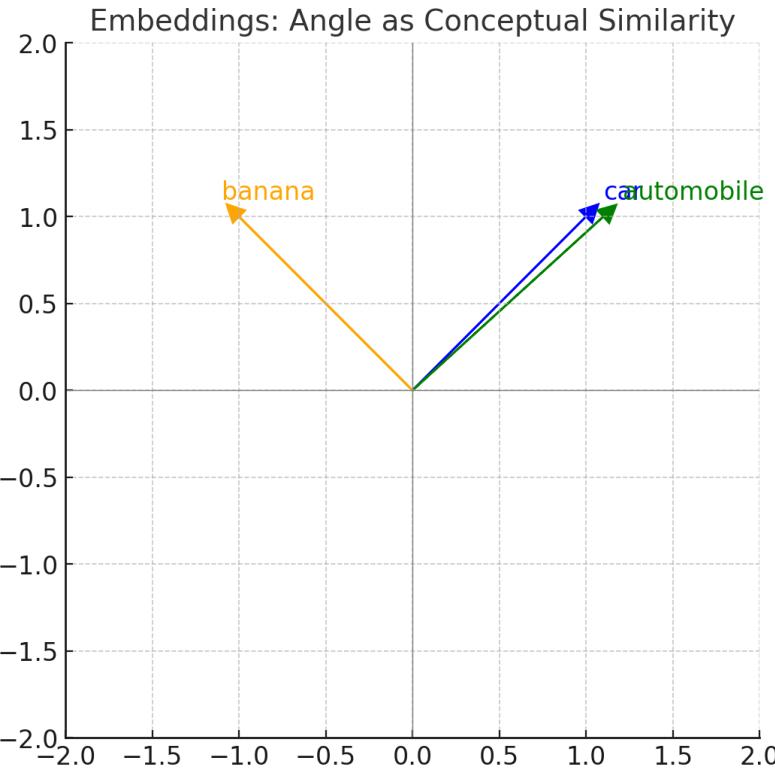
$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$



# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



$$A = [2, 3], \quad B = [4, -1]$$

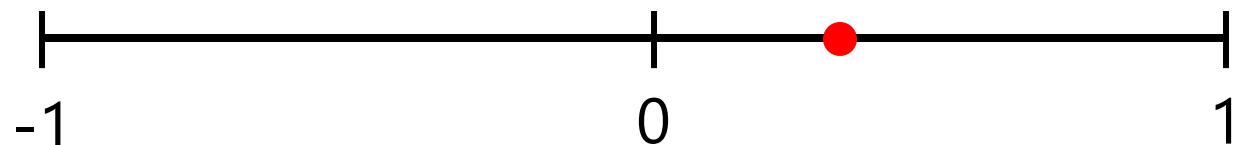
$$A \cdot B = (2 \times 4) + (3 \times -1) = 8 - 3 = 5$$

$$\|A\| = \sqrt{2^2 + 3^2} = \sqrt{4 + 9} = \sqrt{13} \approx 3.606$$

$$\|B\| = \sqrt{4^2 + (-1)^2} = \sqrt{16 + 1} = \sqrt{17} \approx 4.123$$

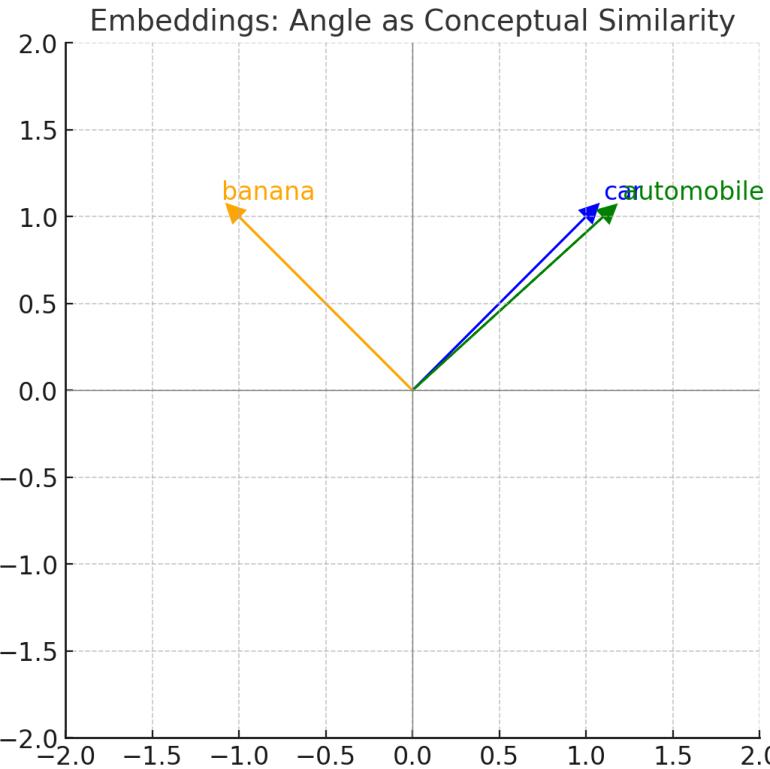
$$\|A\| \times \|B\| \approx 3.606 \times 4.123 \approx 14.85$$

$$\cos(\theta) = \frac{5}{14.85} \approx 0.34$$



# Cosine Similarity

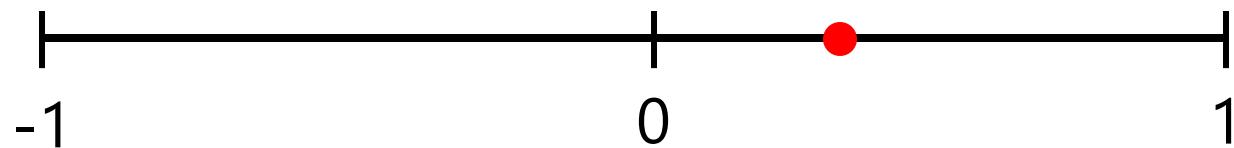
$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



**$\cos(0^\circ) = 1 \rightarrow$  perfect match**

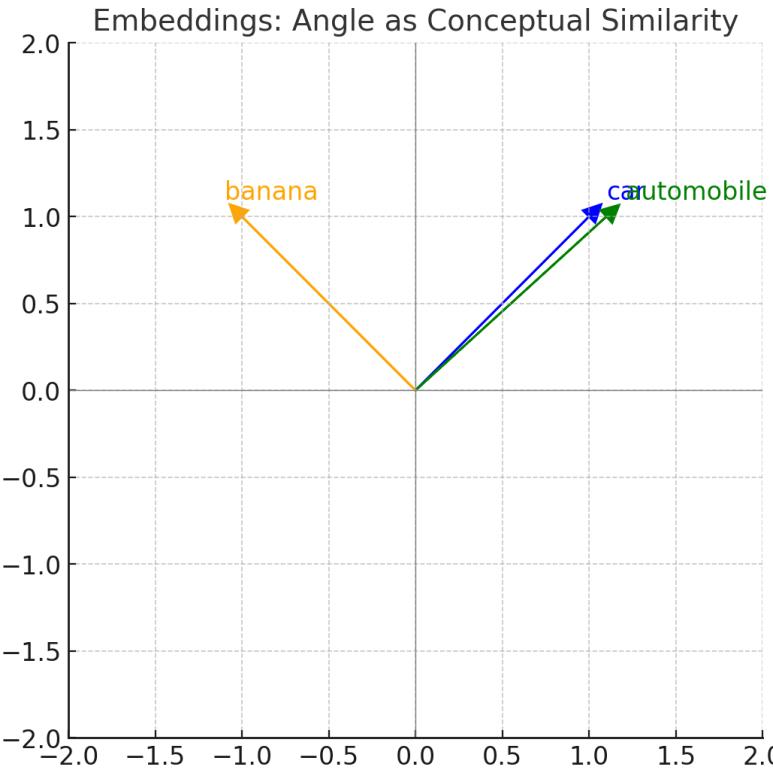
**$\cos(90^\circ) = 0 \rightarrow$  no relation**

**$\cos(180^\circ) = -1 \rightarrow$  opposite**



# Cosine Similarity

$$\text{cosine similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$



**$\cos(0^\circ) = 1 \rightarrow$  perfect match**

**$\cos(90^\circ) = 0 \rightarrow$  no relation**

**$\cos(180^\circ) = -1 \rightarrow$  opposite**

"car" → vector A

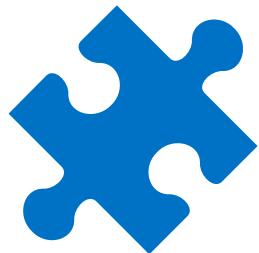
"automobile" → vector B

Their angle is tiny → high similarity

"car" vs. "banana" → angle ~90° → not related.

# Semantic Search Tools in the .NET Ecosystem

# Frameworks & Libraries



Microsoft.Extensions.AI

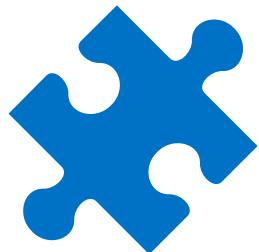


Semantic Kernel



ML.NET

# Frameworks & Libraries

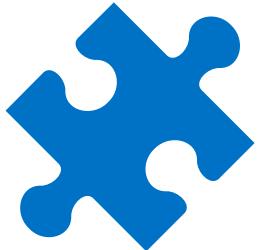


Microsoft.Extensions.AI



ML.NET

# Frameworks & Libraries

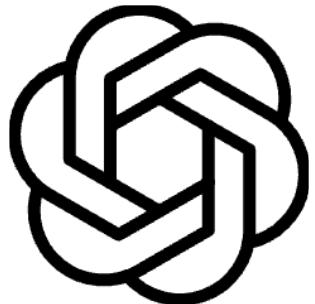


Microsoft.Extensions.AI



ML.NET

# Embedding Models



OpenAI API



Ollama



Hugging Face

# Embedding Models



OpenAI API



★ Ollama

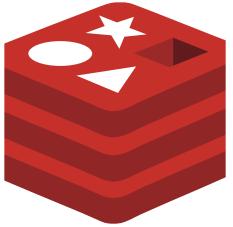


Hugging Face

# Vector Databases



Cosmos DB



Redis



Qdrant



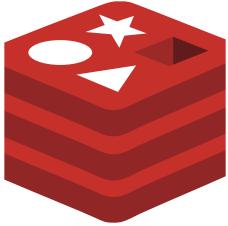
Pinecone,  
Weaviate, Milvus

NOTE: SQL Server 2025 includes a vector data type

# Vector Databases



Cosmos DB



Redis



★ Qdrant



Pinecone,  
Weaviate, Milvus

NOTE: SQL Server 2025 includes a vector data type

# Cloud Services



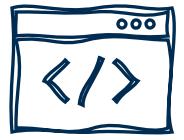
Azure AI Search



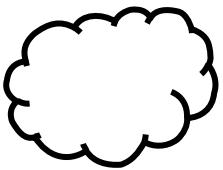
ElasticSearch

# Implementing Semantic Search in .NET

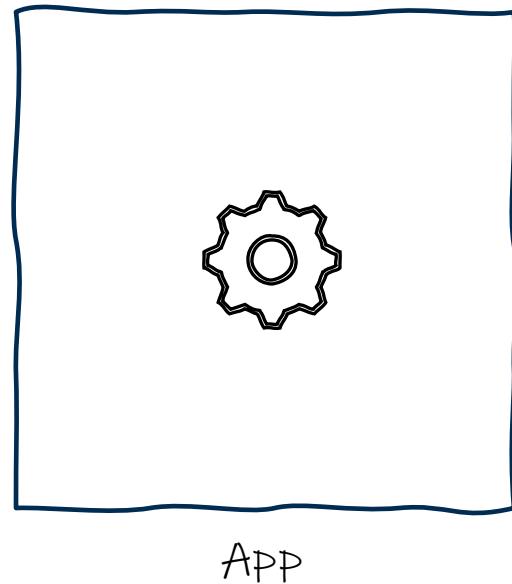
# Semantic Search Process



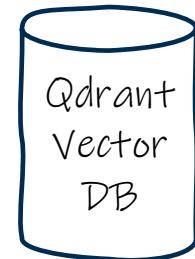
Trailhead  
RSS Feed



Internet

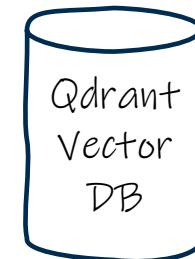
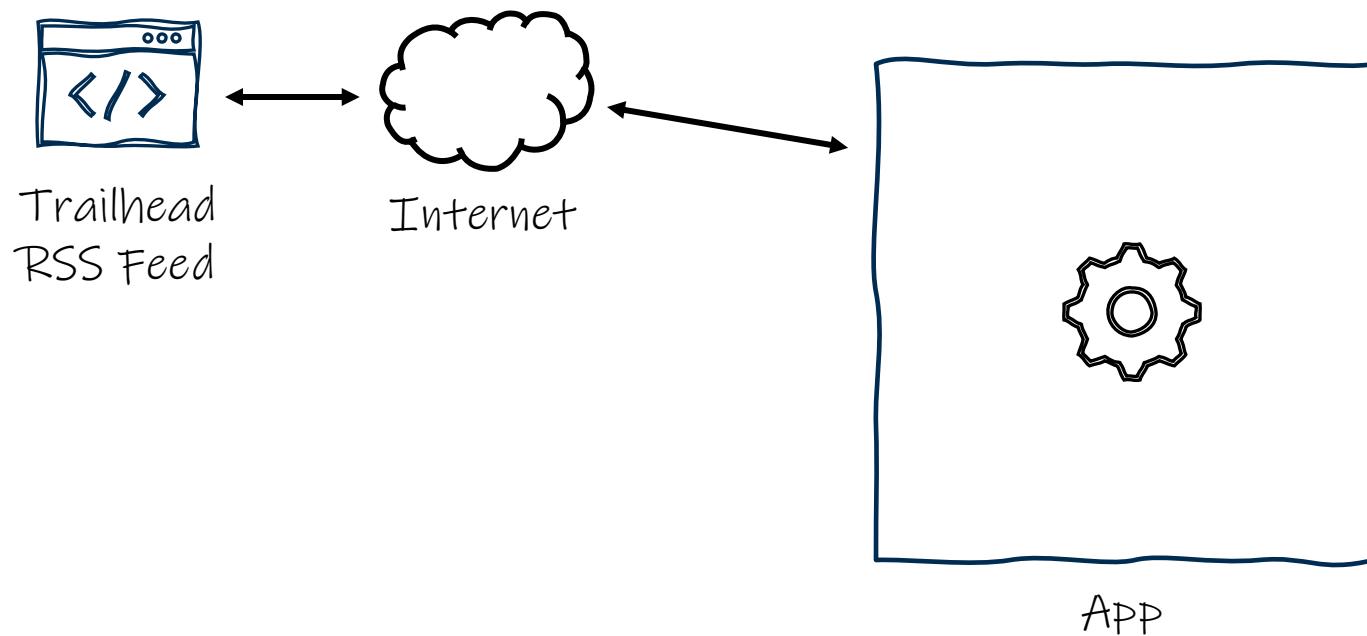


Ollama  
Embedding  
Model

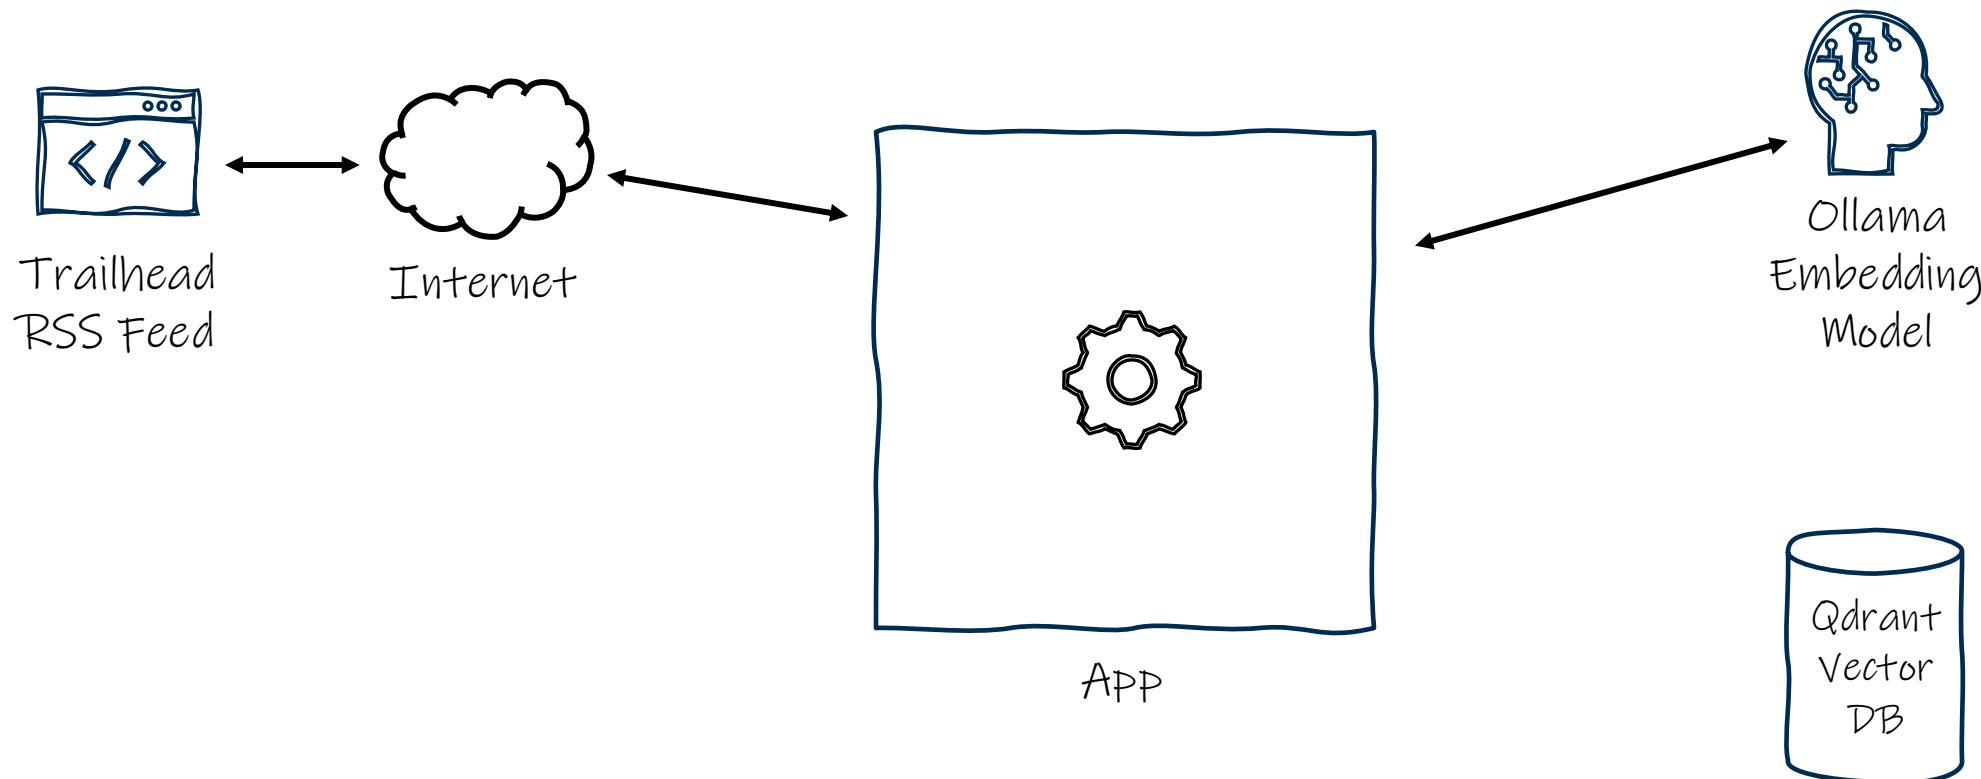


Qdrant  
Vector  
DB

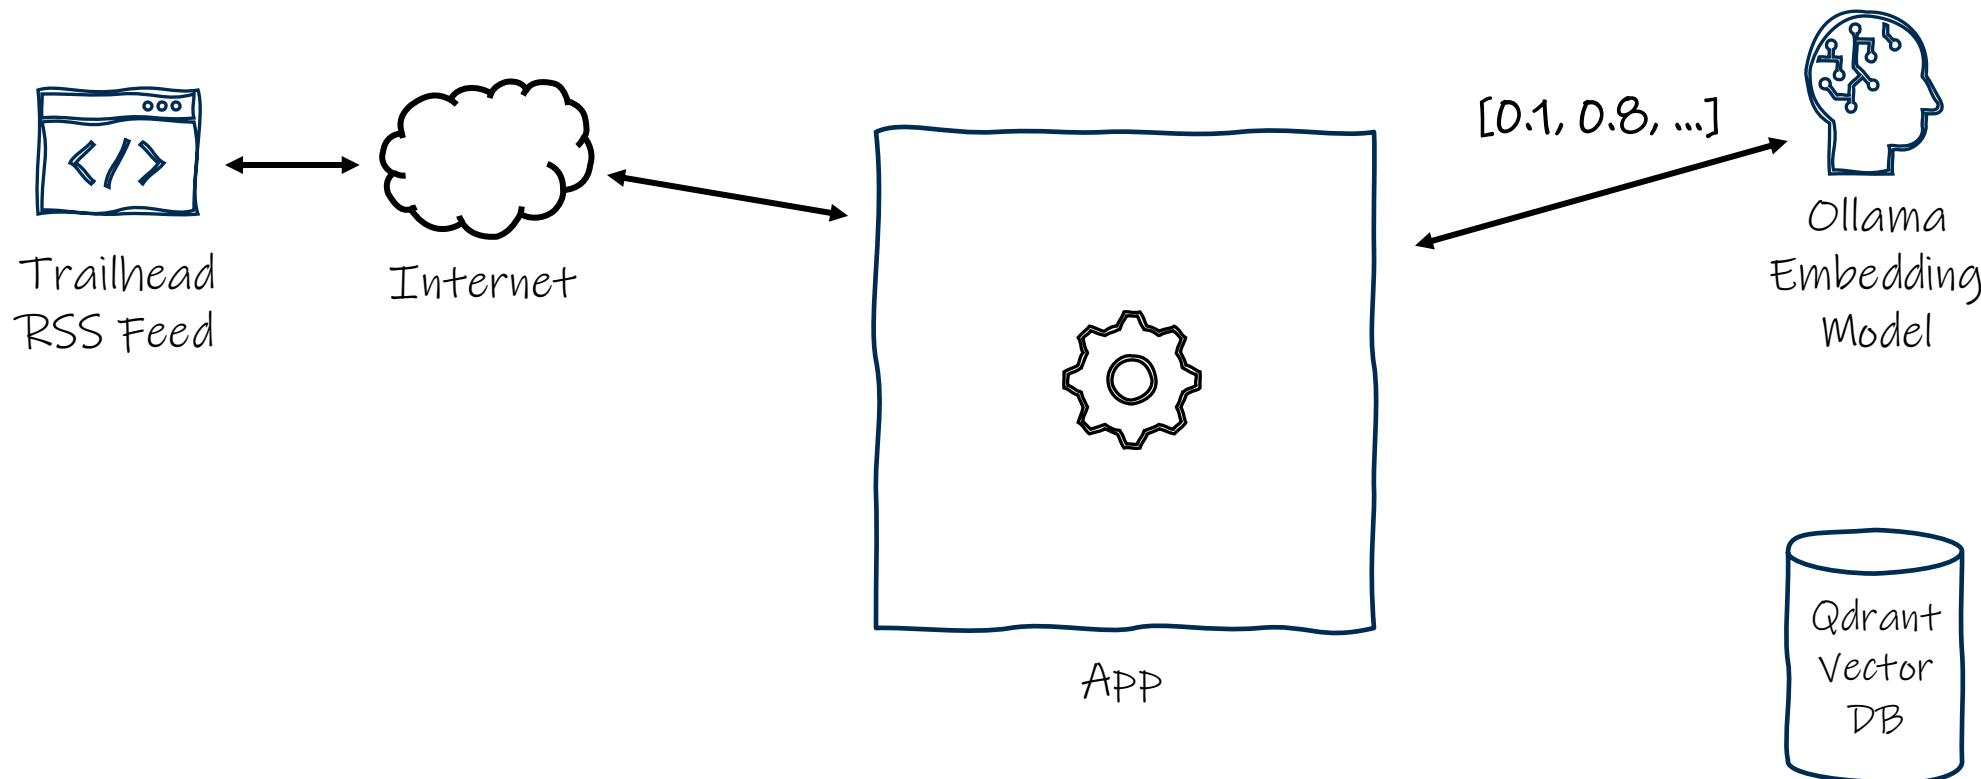
# Semantic Search Process



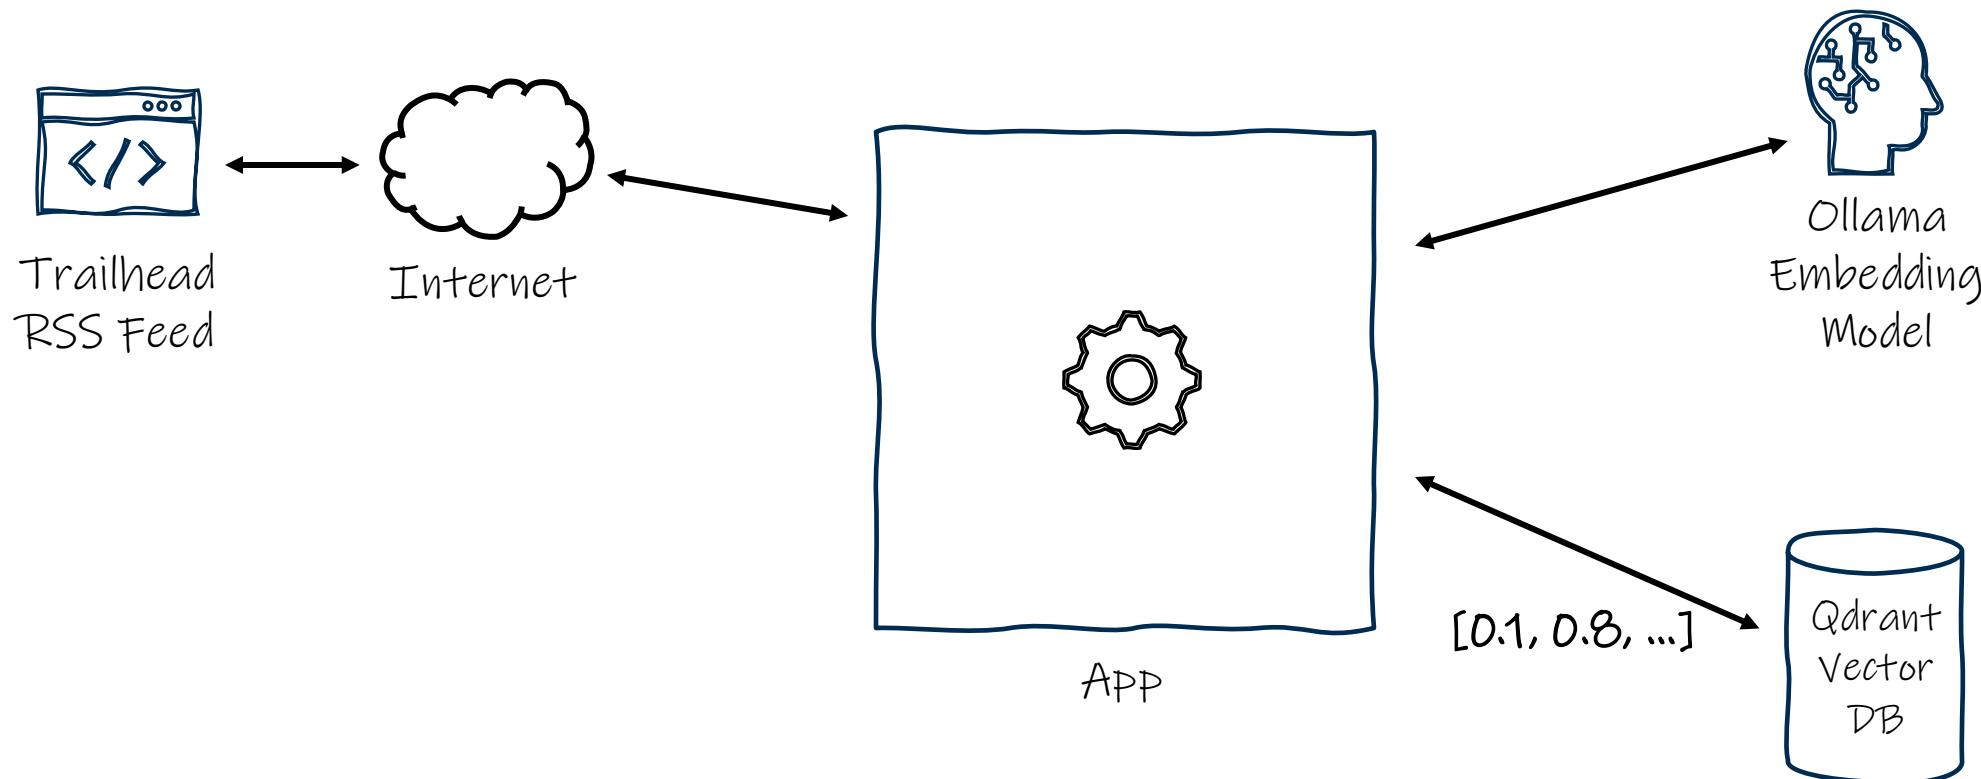
# Semantic Search Process



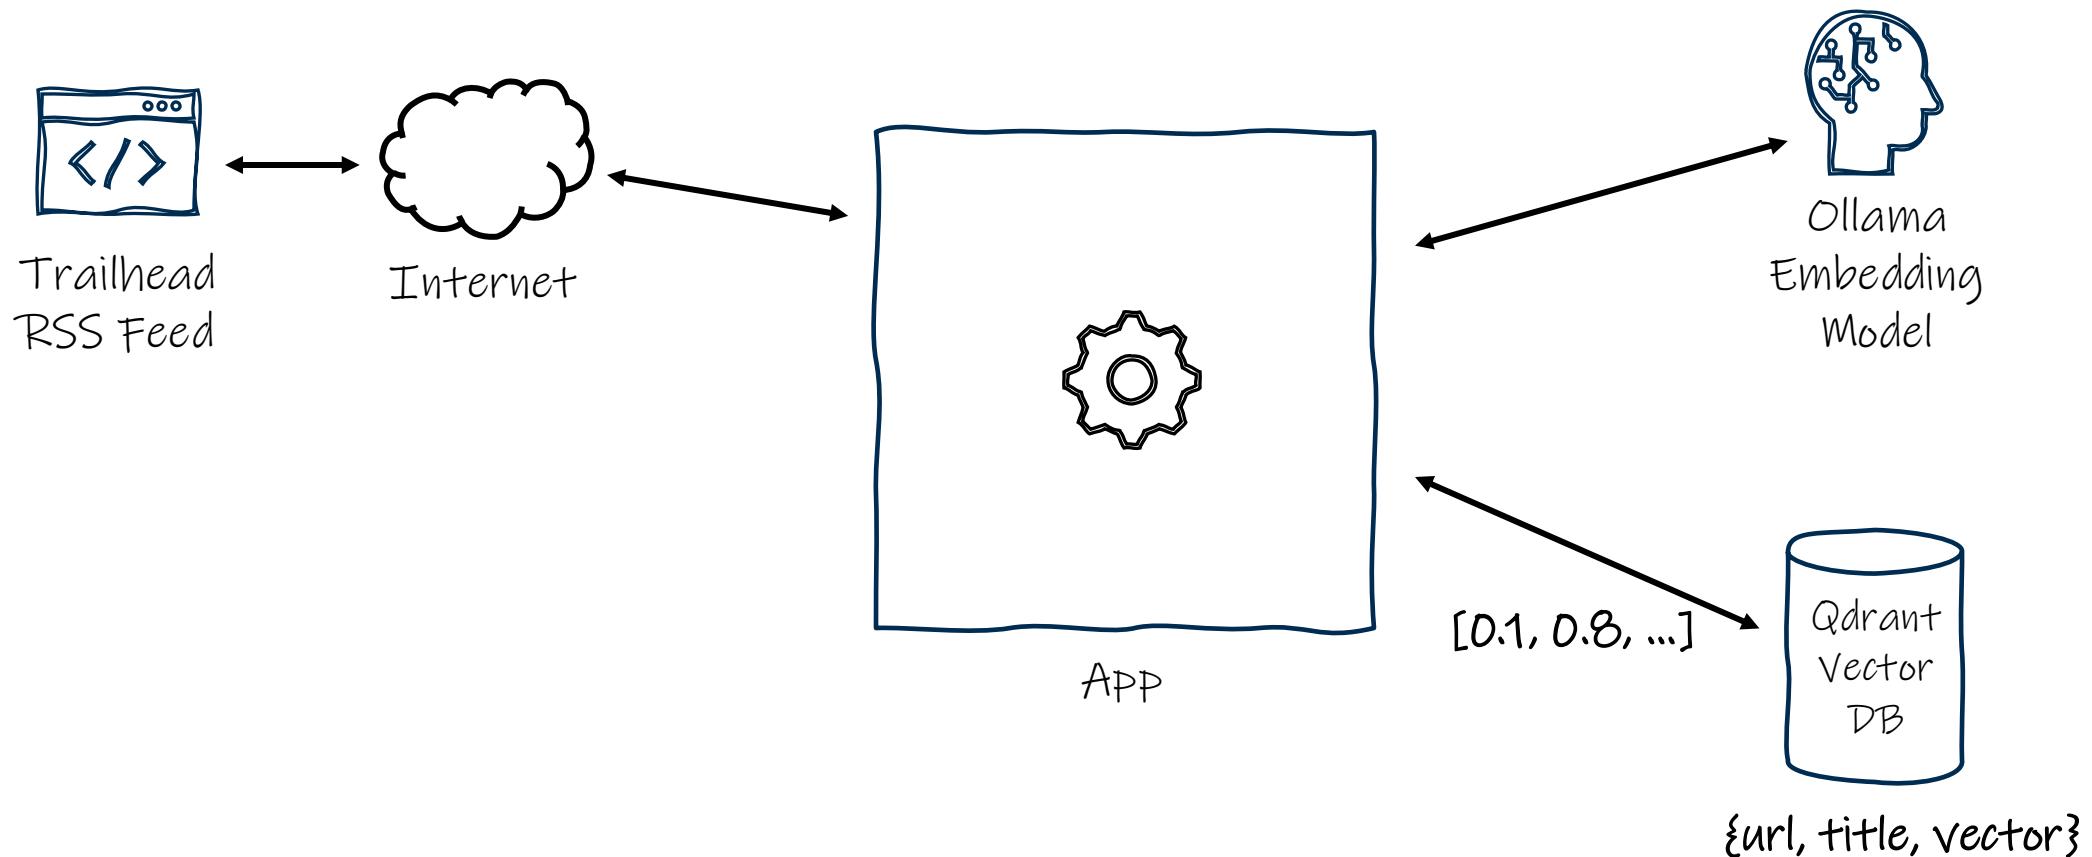
# Semantic Search Process



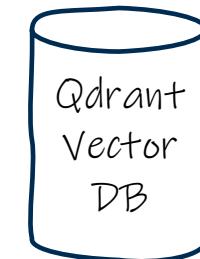
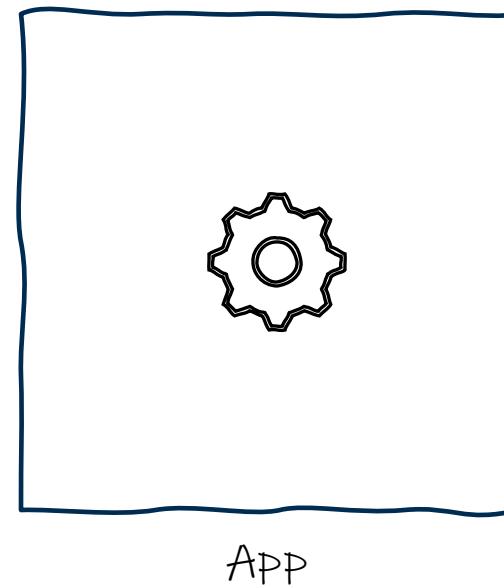
# Semantic Search Process



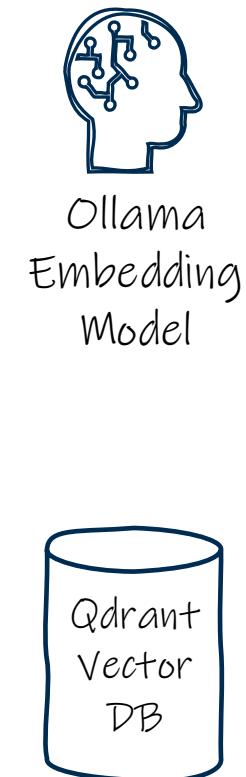
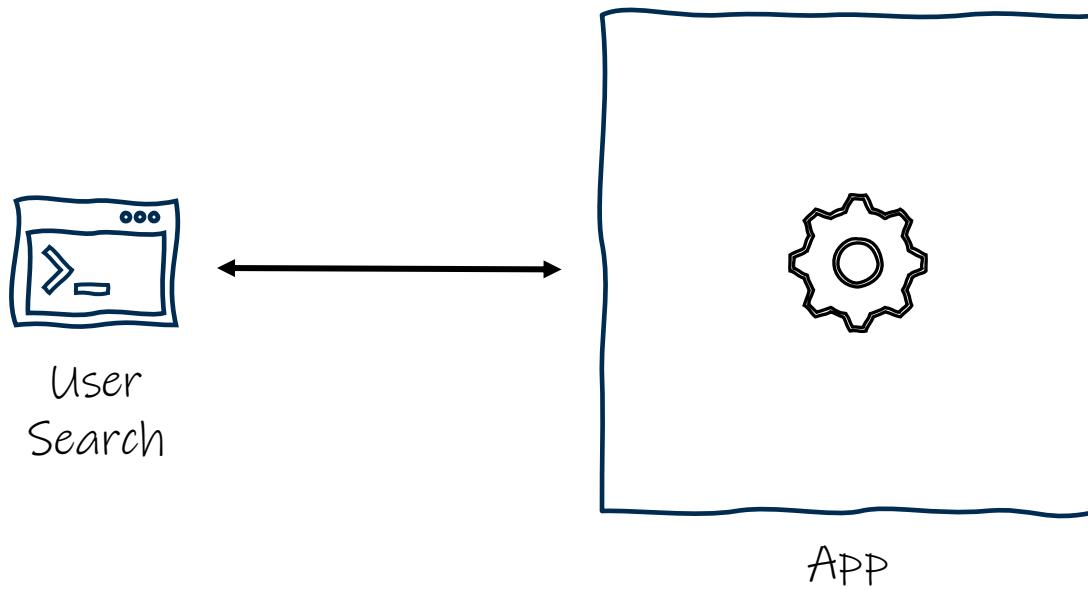
# Semantic Search Process



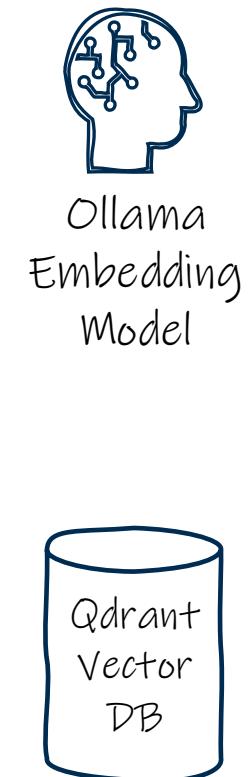
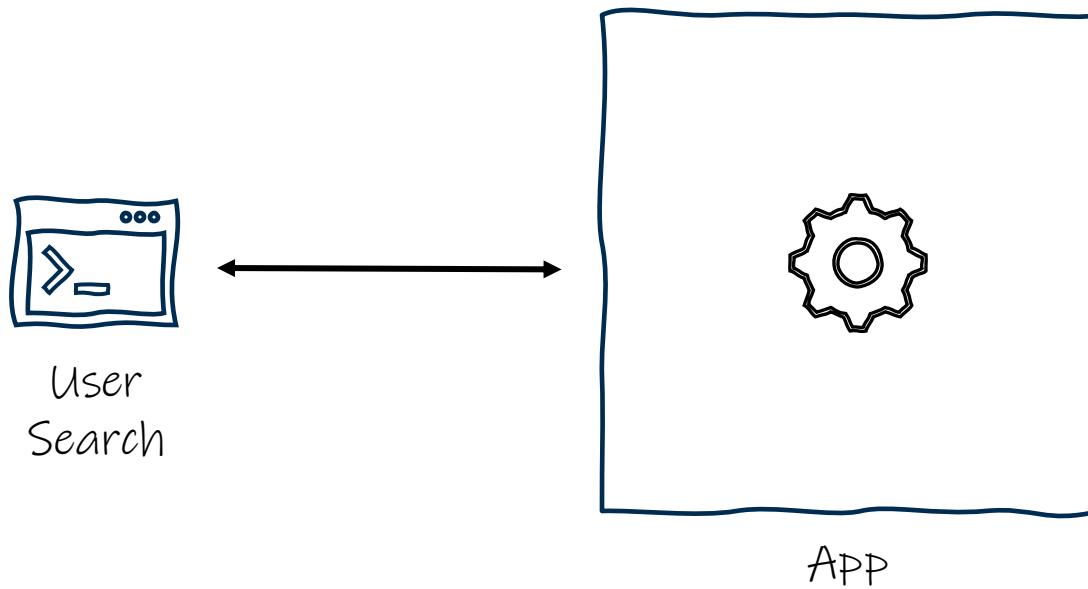
# Semantic Search Process



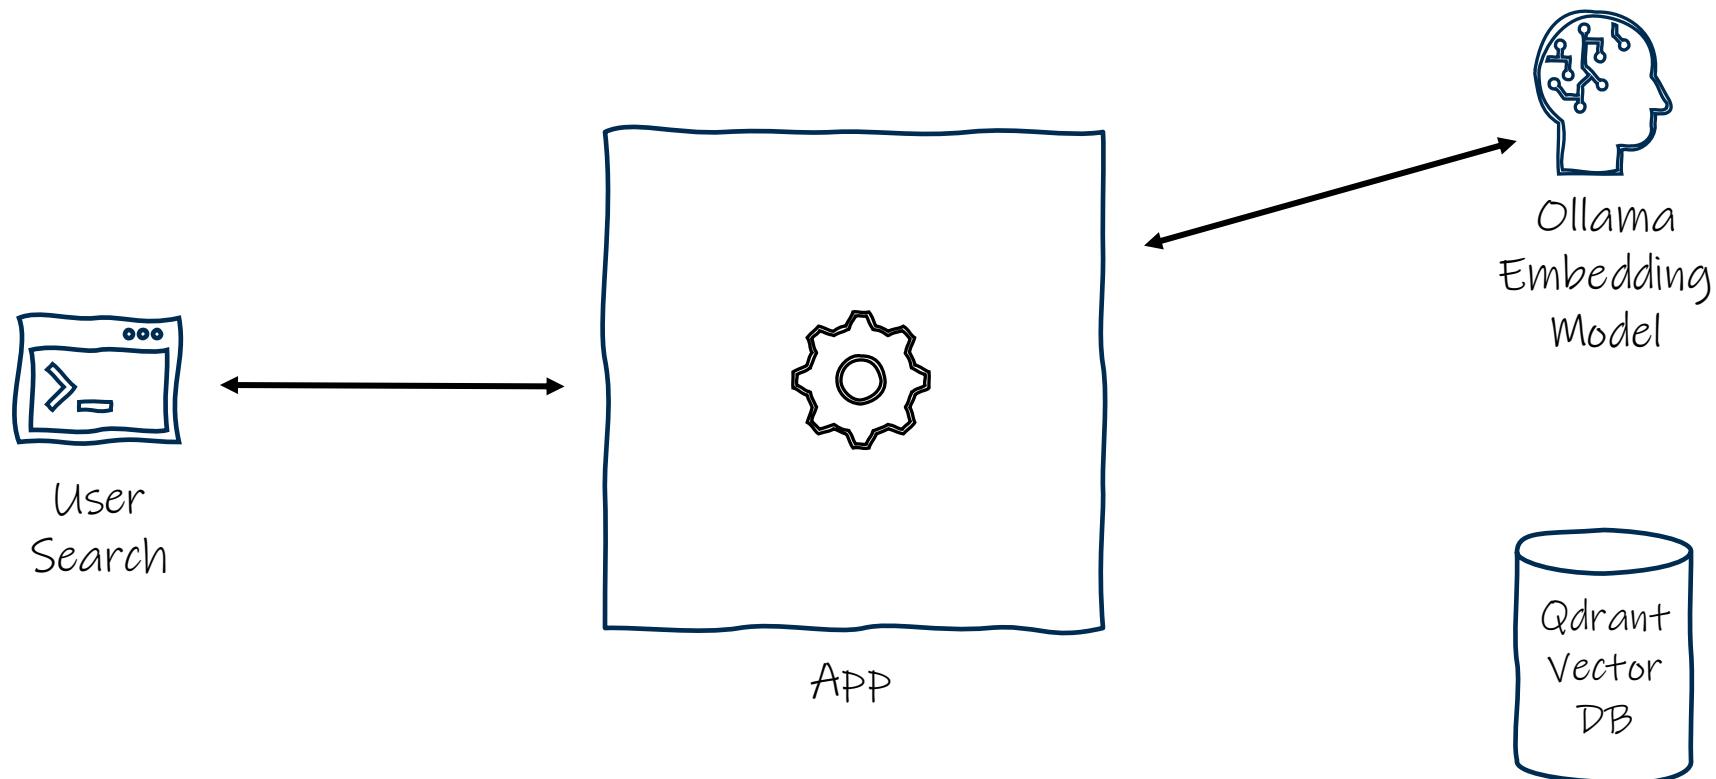
# Semantic Search Process



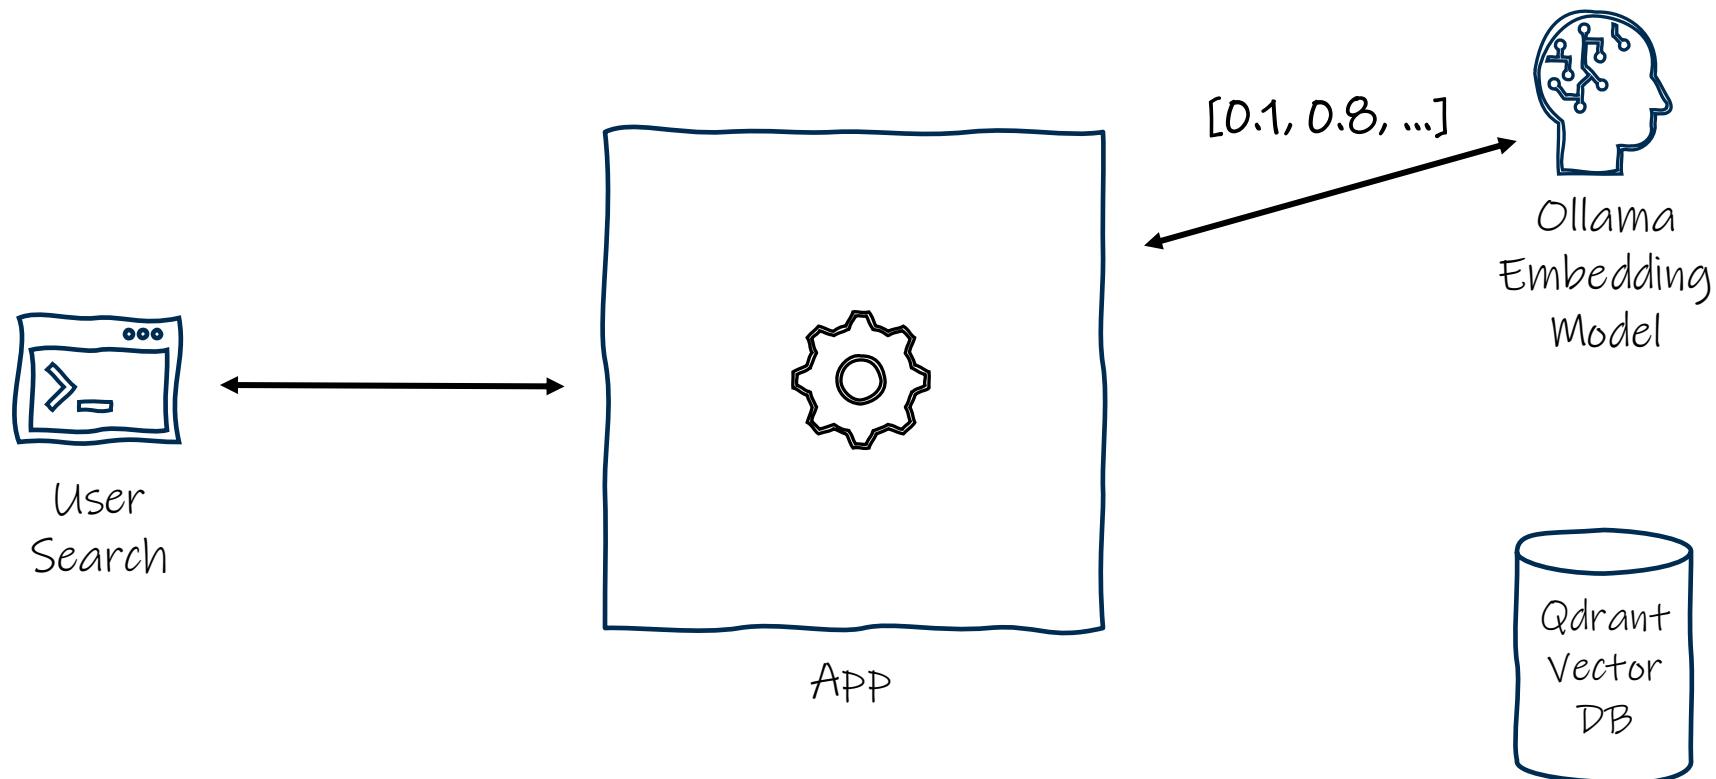
# Semantic Search Process



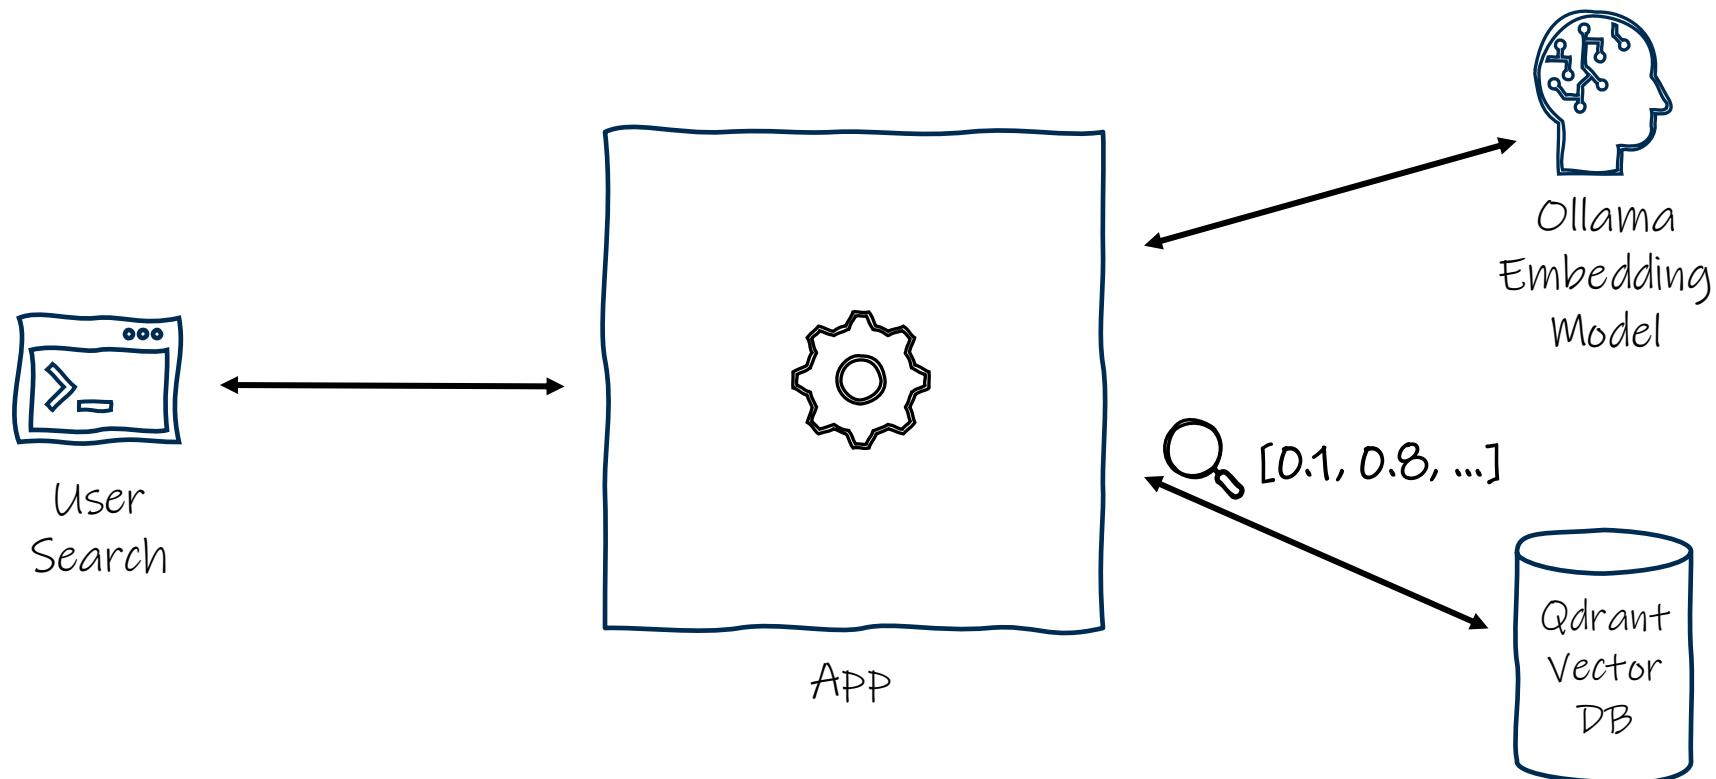
# Semantic Search Process



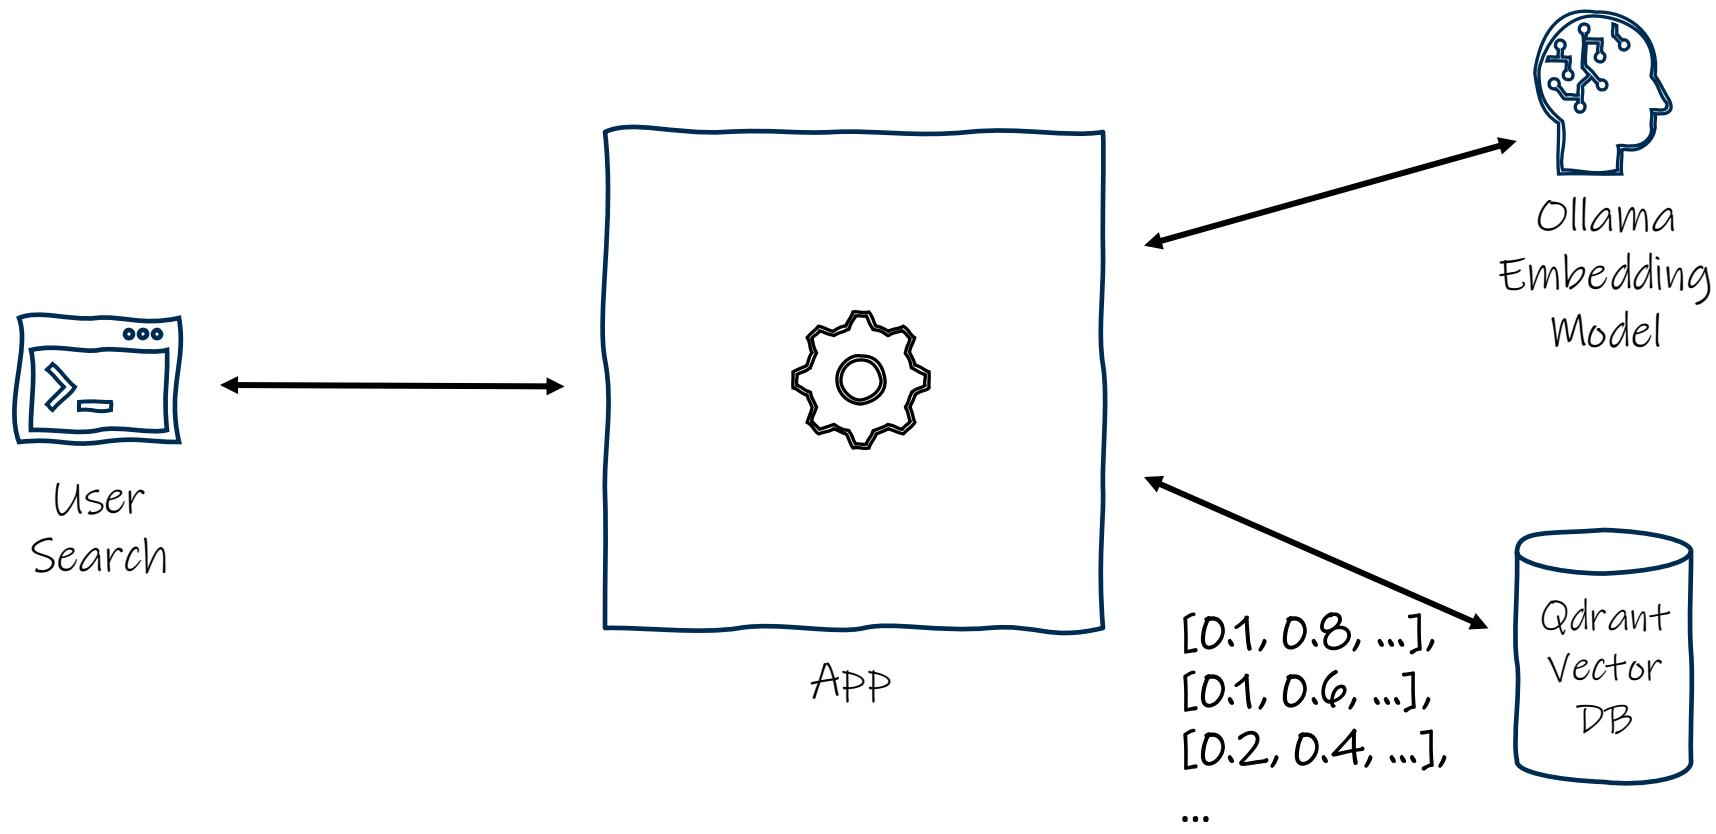
# Semantic Search Process



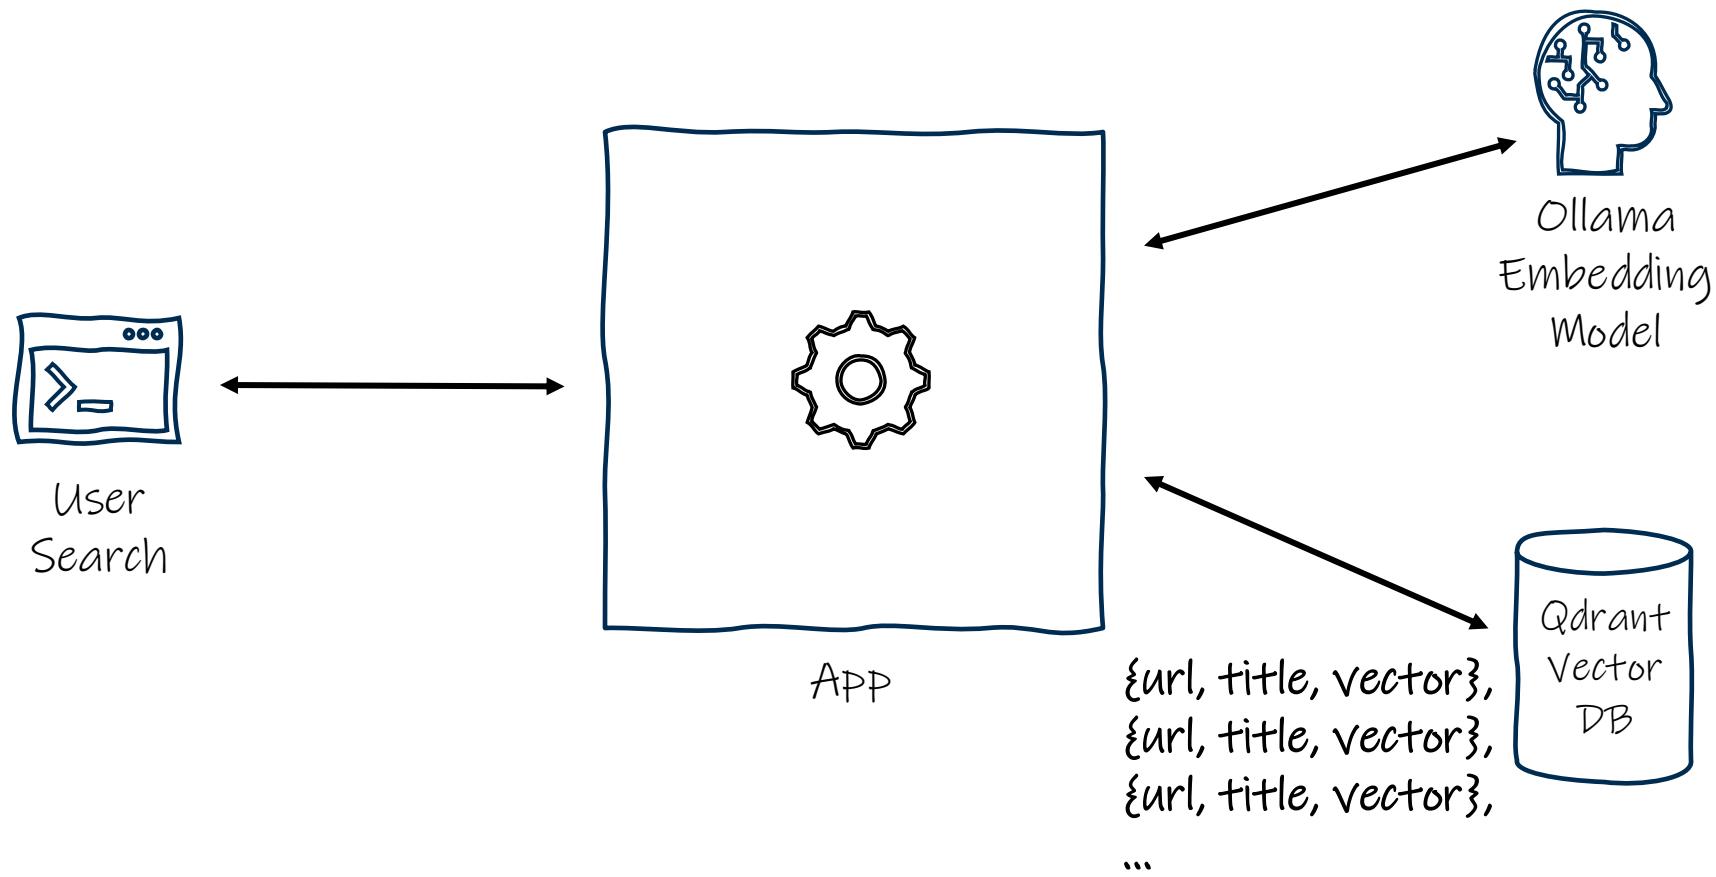
# Semantic Search Process



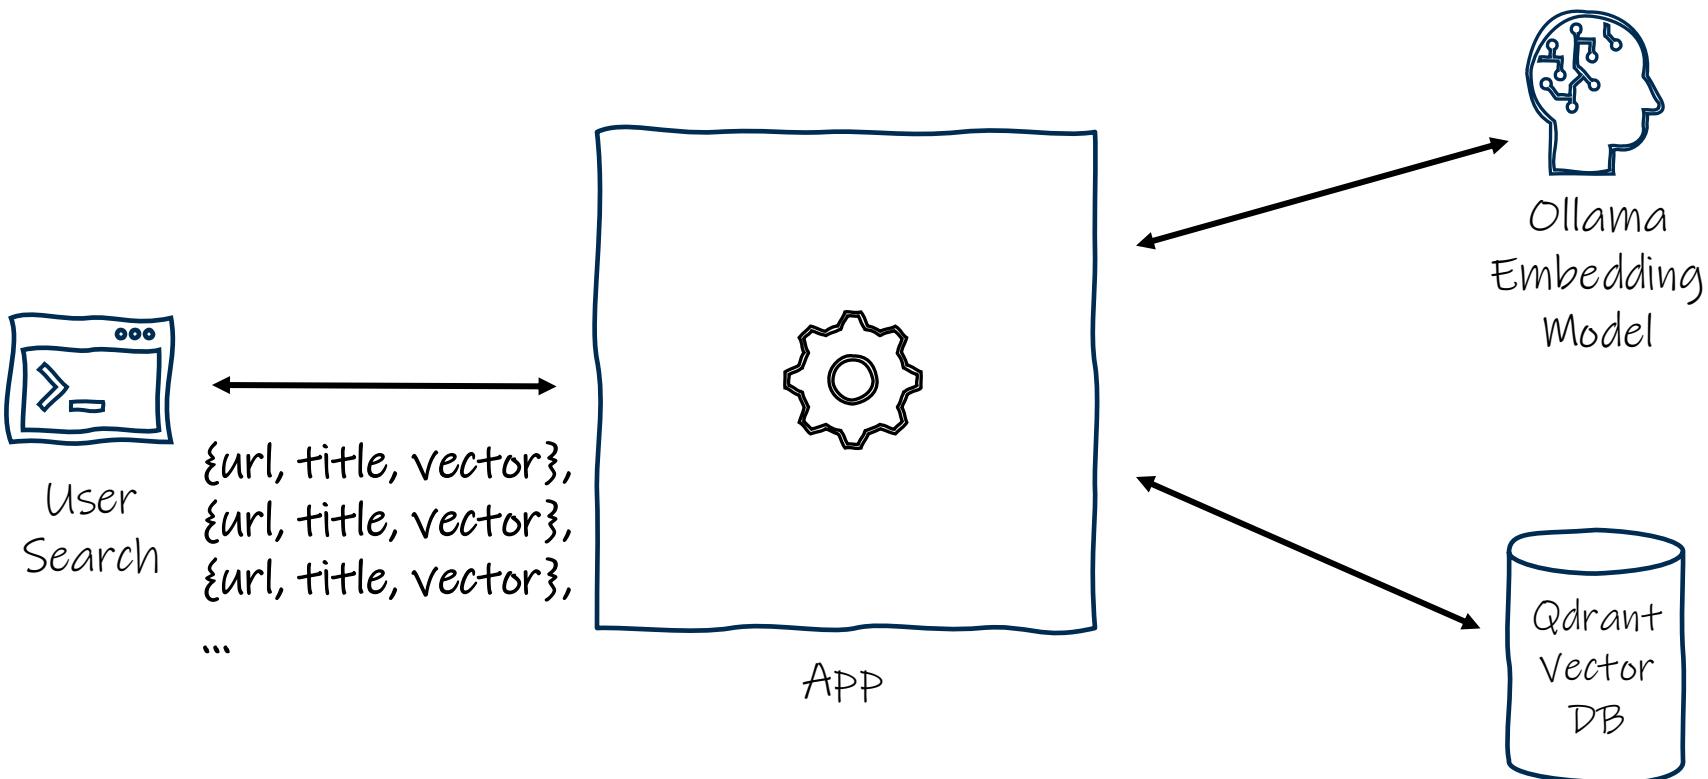
# Semantic Search Process



# Semantic Search Process



# Semantic Search Process



# LIVE DEMO

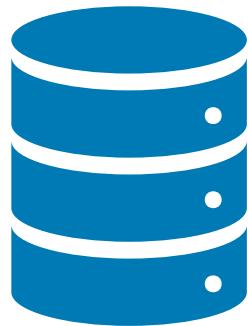


# Practical Considerations

# Cost & Latency Trade-Offs

Option	Cost	Latency	Accuracy
<b>Local Models</b>	free (typically)	medium	medium
<b>Cloud Models</b>	higher	low	high
<b>Hybrid</b>	balanced	balanced	balanced

# Scalability & Storage



Store in a Vector DB



Index Vectors

# Scalability & Storage



1 vector = 1 KB

# Scalability & Storage



1 vector = 1 KB



1M vectors = GBs

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	<ul style="list-style-type: none"><li><span>✓</span> Fast</li><li><span>✓</span> Cheap</li><li><span>✓</span> Lower storage</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> Less nuance</li><li><span>⚠</span> Lower accuracy</li></ul>	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	<ul style="list-style-type: none"><li><span>✓</span> Higher accuracy</li><li><span>✓</span> Captures subtle meaning</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> More compute</li><li><span>⚠</span> Higher storage/cost</li></ul>	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	<ul style="list-style-type: none"><li><span>✓</span> Tuned for specific language (legal, medical, finance, etc.)</li><li><span>✓</span> Often best results</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> May miss general queries</li><li><span>⚠</span> Limited availability</li></ul>	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
<b>Small embeddings (384–768 dims)</b>	Fast Cheap Lower storage	Less nuance Lower accuracy	Quick search, lightweight apps, prototyping
<b>Large embeddings (1024–3000 dims)</b>	Higher accuracy Captures subtle meaning	More compute Higher storage/cost	Production search, nuanced queries, RAG
<b>Domain-specific models</b>	Tuned for specific language (legal, medical, finance, etc.) Often best results	May miss general queries Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
Small embeddings (384–768 dims)	<span>✓</span> Fast <span>✓</span> Cheap <span>✓</span> Lower storage	<span>⚠</span> Less nuance <span>⚠</span> Lower accuracy	Quick search, lightweight apps, prototyping
Large embeddings (1024–3000 dims)	<span>✓</span> Higher accuracy <span>✓</span> Captures subtle meaning	<span>⚠</span> More compute <span>⚠</span> Higher storage/cost	Production search, nuanced queries, RAG
Domain-specific models	<span>✓</span> Tuned for specific language (legal, medical, finance, etc.) <span>✓</span> Often best results	<span>⚠</span> May miss general queries <span>⚠</span> Limited availability	Specialized industries, enterprise apps

# Quality & Model Choice

Type	Pros	Cons	Use Cases
Small embeddings (384–768 dims)	<ul style="list-style-type: none"><li><span>✓</span> Fast</li><li><span>✓</span> Cheap</li><li><span>✓</span> Lower storage</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> Less nuance</li><li><span>⚠</span> Lower accuracy</li></ul>	Quick search, lightweight apps, prototyping
Large embeddings (1024–3000 dims)	<ul style="list-style-type: none"><li><span>✓</span> Higher accuracy</li><li><span>✓</span> Captures subtle meaning</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> More compute</li><li><span>⚠</span> Higher storage/cost</li></ul>	Production search, nuanced queries, RAG
Domain-specific models	<ul style="list-style-type: none"><li><span>✓</span> Tuned for specific language (legal, medical, finance, etc.)</li><li><span>✓</span> Often best results</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> May miss general queries</li><li><span>⚠</span> Limited availability</li></ul>	Specialized industries, enterprise apps

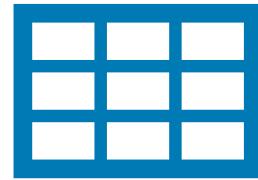
# Quality & Model Choice

Type	Pros	Cons	Use Cases
Small embeddings (384–768 dims)	<ul style="list-style-type: none"><li><span>✓</span> Fast</li><li><span>✓</span> Cheap</li><li><span>✓</span> Lower storage</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> Less nuance</li><li><span>⚠</span> Lower accuracy</li></ul>	Quick search, lightweight apps, prototyping
Large embeddings (1024–3000 dims)	<ul style="list-style-type: none"><li><span>✓</span> Higher accuracy</li><li><span>✓</span> Captures subtle meaning</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> More compute</li><li><span>⚠</span> Higher storage/cost</li></ul>	Production search, nuanced queries, RAG
Domain-specific models	<ul style="list-style-type: none"><li><span>✓</span> Tuned for specific language (legal, medical, finance, etc.)</li><li><span>✓</span> Often best results</li></ul>	<ul style="list-style-type: none"><li><span>⚠</span> May miss general queries</li><li><span>⚠</span> Limited availability</li></ul>	Specialized industries, enterprise apps

# When NOT to Use Semantic Search



Tiny Datasets



Structured Lookups

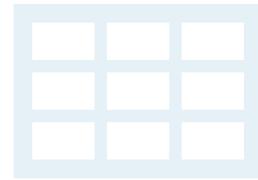


Strict Regulatory  
environments

# When NOT to Use Semantic Search



Tiny Datasets



Structured Lookups

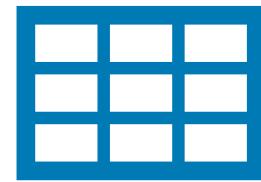


Strict Regulatory  
environments

# When NOT to Use Semantic Search



Tiny Datasets



Structured Lookups

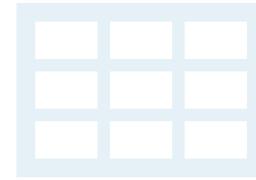


Strict Regulatory  
environments

# When NOT to Use Semantic Search



Tiny Datasets



Structured Lookups



Strict Regulatory  
environments

# Summing Up

1. Semantic search **searches meaning**, not just words or parts of words
2. Powered by **vectors** and **embeddings**
3. Many **tools exist** such MEAI, Ollama, Azure OpenAI API, Qdrant, etc.
4. Balance tradeoffs of **local** vs **hosted** models.



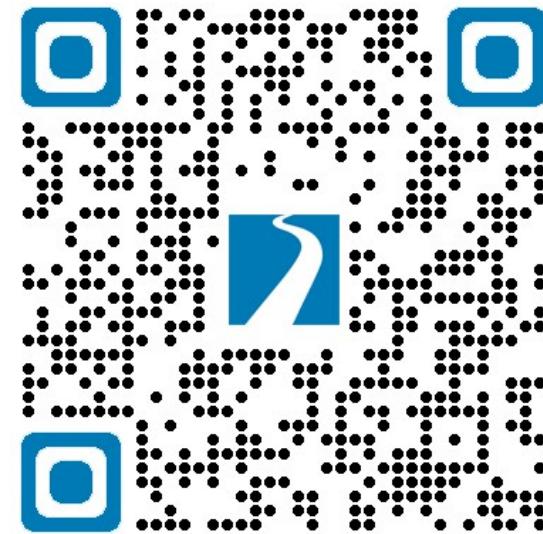
# Thanks! Questions?

## Jonathan "J." Tower

- ✉ Microsoft MVP in .NET
- ✉ .NET Foundation Board of Directors
- ✉ jtower@trailheadtechnology.com
- ✉ trailheadtechnology.com/blog
- ✉ jtowermi
- ✉ Jonathan "J." Tower

[github.com/trailheadtechnology/api-security](https://github.com/trailheadtechnology/api-security)

**EXPERT**  
Consultation



bit.ly/th-offer